

Towards Robust and Secure Embodied AI: A Survey on Vulnerabilities and Attacks

WENPENG XING, Zhejiang University, Hangzhou, China and Binjiang Institute of Zhejiang University, Hangzhou, China

MINGHAO LI, College of Big Data and Software Engineering, Chongqing University, Chongqing, China

MOHAN LI, Cyberspace Institute of Advanced Technology, Guangzhou University, Guangzhou, China

MENG HAN, Zhejiang University, Hangzhou, China, Binjiang Institute of Zhejiang University, Hangzhou, China, and GenTel.io, hangzhou, China

Embodied AI systems, integrating Large Vision-Language Models (LVLMs) and Large Language Models (LLMs) with physical actuators and sensors, face unique robustness and security challenges stemming from the complex interplay between perception, cognition, and actuation in real-world environments. This survey provides a systematic analysis of these vulnerabilities and associated attack surfaces. We propose a tripartite vulnerability taxonomy comprising foundational, integration, and contextual risks. Foundational vulnerabilities arise from inherent limitations in current AI architectures and training paradigms; Integration vulnerabilities emerge from the composition of cyber-physical components; And contextual vulnerabilities stem from dynamic physical environments and deployment conditions. Correspondingly, we present a comprehensive attack taxonomy that encompasses foundational attacks on LLMs/LVLMs (including logits-based, optimization-based, prompt-based, and cross-modality attacks), integration-level cybersecurity threats (such as man-in-the-middle, firmware, side-channel, and supply chain attacks), and contextual attacks (primarily sensor spoofing across multiple modalities). We further examine representative failure modes of the cognitive core, review existing evaluation methodologies and benchmarks, and synthesize a multi-layer defense framework that integrates perceptual redundancy, runtime monitoring, and hardware-enforced safety mechanisms. This work offers a unified conceptual framework and practical roadmap for designing and evaluating robust, secure Embodied AI systems in safety-critical real-world deployments.

CCS Concepts: • **General and reference** → **Surveys and overviews**; • **Security and privacy** → *Software and application security*; • **Computer systems organization** → *Embedded and cyber-physical systems*; • **Computing methodologies** → *Machine learning*;

Additional Key Words and Phrases: Large vision language models, large language models, Embodied AI, adversarial attack

This research was supported by National Natural Science Foundation of China under Grant No. 62376246 and Zhejiang Provincial Natural Science Foundation of China under Grant No. LMS26F020027.

Authors' Contact Information: Wenpeng Xing, Zhejiang University, Hangzhou, Zhejiang, China and Binjiang Institute of Zhejiang University, Hangzhou, China; e-mail: wpxing@zju.edu.cn; Minghao Li, College of Big Data and Software Engineering, Chongqing University, Chongqing, Chongqing, China; e-mail: mhli@stu.cqu.edu.cn; Mohan Li, Cyberspace Institute of Advanced Technology, Guangzhou University, Guangzhou, Guangdong, China; e-mail: limohan.hit@gmail.com; Meng Han (corresponding author), Zhejiang University, Hangzhou, Zhejiang, China, Binjiang Institute of Zhejiang University, Hangzhou, China, and GenTel.io, hangzhou, China; e-mail: mhan@zju.edu.cn.



This work is licensed under a Creative Commons Attribution 4.0 International License.

© 2026 Copyright held by the owner/author(s).

ACM 0360-0300/2026/05-ART312

<https://doi.org/10.1145/3806048>

ACM Reference Format:

Wenpeng Xing, Minghao Li, Mohan Li, and Meng Han. 2026. Towards Robust and Secure Embodied AI: A Survey on Vulnerabilities and Attacks. *ACM Comput. Surv.* 58, 12, Article 312 (May 2026), 36 pages. <https://doi.org/10.1145/3806048>

1 Introduction

The rapid advancement of AI has established Embodied AI as a cornerstone of autonomous driving and industrial automation, defined by the synergy of *autonomy*, *embodiment*, and *cognition* [10]. While enabling complex real-world interactions, these core characteristics inherently expand the threat surface (Figure 1): *Autonomy* introduces decision-making fragility in dynamic environments; *Embodiment* exposes physical sensors to jamming and spoofing attacks; And *cognition* renders the algorithmic core vulnerable to adversarial manipulation. Furthermore, deployment in sensitive settings necessitates rigorous data governance compliant with regulations like the *EU AI Act* [31]. Consequently, addressing these vulnerabilities—which threaten to cause physical safety violations, unauthorized control, and privacy breaches—is paramount for ensuring reliable deployment.

The remainder of this article is structured as follows, as illustrated in Figure 2. **Section 2** establishes a tripartite vulnerability taxonomy, categorizing risks into **foundational**, **integration**, and **contextual** vulnerabilities. **Section 3** analyzes the corresponding attack landscape, spanning **threat models**, **foundational** attacks on LLMs/LVLMs, **integration** attacks (cybersecurity threats), and **contextual** attacks (sensor spoofing). **Section 4** examines the challenges and failure modes of the AI core components, including **cognitive**, **security**, and **operational** failures. **Section 5** reviews **evaluation** methodologies and **benchmarks** for assessing the safety and robustness of Embodied AI systems. **Section 6** synthesizes a **multi-layer defense framework** for building robust Embodied AI, covering perceptual, cognitive, and physical safety layers. Finally, the conclusion is presented in **Section 7**.

2 Vulnerability Taxonomy

Embodied AI vulnerabilities arise from three complementary aspects: **foundational** limitations in AI methodologies, **integration** challenges in cyber-physical systems, and **contextual** dependencies on deployment environments. This taxonomy organizes vulnerabilities by their nature rather than their manifestation, providing clearer guidance for targeted defense strategies.

2.1 Foundational Vulnerabilities

Foundational vulnerabilities originate from the core design choices and limitations of contemporary AI approaches:

- **Training Paradigm Misalignment:** The autoregressive (next-word prediction) objective prioritizes statistical coherence over safety, truthfulness, or harmlessness. Safety fine-tuning offers limited mitigation against adversaries [11].
- **Data-Induced Biases and Artifacts:** Pre-training on uncensored internet data introduces systematic biases [28] and toxic content, creating predictable failure modes that adversaries can reliably trigger.
- **Architectural Susceptibility to Gradient-Based Attacks:** Deep Neural Networks’ near-linear behavior in high-dimensional spaces [113] enables efficient gradient exploitation via methods like *Fast Gradient Sign Method (FGSM)* [40] and *Carlini-Wagner (CW)* attacks [12]. Incomplete input space coverage during training creates “blind spots” [90, 92] that adversaries systematically discover.




Fig. 1. Security Threats and Challenges in Embodied AI Systems. A six-dimensional taxonomy of security threats in Embodied AI systems: (a) environmental challenges affecting perception; (b) sensor-specific jamming/spoofing attacks; (c) diverse attack vectors from firmware to supply chain; (d) internal failures including mechanical and software vulnerabilities; (e) functional vulnerabilities encompassing ethical and robustness concerns; and (f) severe real-world consequences resulting from these attacks.

2.2 Integration Vulnerabilities

Integration vulnerabilities manifest when AI components are assembled into complete **cyber-physical** systems, creating risks not present in isolated components:

- **Multimodal Attack Surface Expansion:** LVLs create novel vulnerabilities where adversarial perturbations in one modality (e.g., manipulated visuals) propagate to affect reasoning in another (e.g., textual decisions) [72].
- **Text-to-Action Representation Gaps:** Unlike traditional software jailbreaks, attackers can target the structured outputs (JSON, YAML) that bridge high-level planning and low-level control [118], enabling direct manipulation of physical behavior through the system’s own action representation formats.
- **Simulation-Reality Discrepancies in Generalist Models:** Models like *OpenVLA* [59] and π_0 [7] enable cross-embodiment generalization but propagate physical sensor noise through control loops [64], creating distinct safety assurance challenges across simulation and reality.

2.3 Contextual Vulnerabilities

Contextual vulnerabilities vary with deployment **environment**, **adversary** capabilities, and **operational** conditions:

- **Environmental Dynamics and Sensor Limitations:** Real-world conditions (lighting changes, weather, partial obstructions) induce sensor misinterpretations [125] with severe consequences, as evidenced by autonomous vehicle accidents [121]. The reliance on sensor data [47] creates systemic exposure to environmental variations.
- **Hardware-Software Co-Dependencies and Degradation:** Mechanical wear in industrial and medical robots leads to operational disruptions [60], while software bugs and edge-case oversights create failure modes that depend on operational duration and specific conditions.




Fig. 2. Overview of the Survey Structure. The hierarchical organization shows the logical flow from vulnerabilities to attacks, failure modes, evaluation, and defense strategies.

- **Human-System Interaction Failures:** Collaborative systems (cobots, drones) expose vulnerabilities at human-machine interfaces where safety protocol breaches—exacerbated by human error or malicious intent—can lead to catastrophic failures. Natural language misinterpretation in systems like *Talk2Car* [27] and *ALFRED* [109] similarly depends on specific instruction contexts.
- **Cyber-Physical Attack Convergence:** Integration with **Internet of Things (IoT)** and **Cyber-Physical Systems (CPS)** expands attack surfaces where physical tampering [30] converges with




Fig. 3. Comprehensive Security Architecture for Embodied AI Systems: Mapping the three-tier vulnerability taxonomy to system components and defensive safeguards.

cyber exploits [75, 98], enabling novel threat vectors like the Mirai botnet DDoS (Distributed Denial of Service) attacks and **Global Positioning System (GPS)** spoofing [51, 132].

Summary: Vulnerability Taxonomy for Embodied AI

This section establishes a three-layer vulnerability taxonomy for Embodied AI systems, analyzing how AI limitations translate into real-world risks. The taxonomy comprises:

- **Foundational Vulnerabilities:** Rooted in inherent AI limitations, including training misalignment, data biases, and architectural susceptibility to adversarial attacks.
- **Integration Vulnerabilities:** Emerging from system composition, such as multimodal attack surfaces, text-to-action gaps, and simulation-reality discrepancies.
- **Contextual Vulnerabilities:** Dependent on deployment environments, including environmental dynamics, hardware-software co-dependencies, and human-system interaction failures.

3 Attack on Vulnerabilities

Embodied AI systems inherit vulnerabilities from both their AI core components and their cyber-physical integration. To systematically analyze the corresponding attack landscape, we adopt the three-layer vulnerability taxonomy established in Section 2. As visually summarized in Figure 3, attacks not only target distinct architectural layers but can also propagate sequentially across the **perception-planning-control** pipeline, which directly motivates our analytical framework.

We begin by formalizing a **Threat Model** (Section 3.1) that defines attacker capabilities and objectives across these layers. The subsequent analysis is then structured accordingly: we present a three-layer **Attack Taxonomy** (Section 3.2) to systematically analyze the vulnerabilities of Embodied AI in foundational modules, component integration, and physical context.

3.1 Threat Model

An adversarial threat model is defined by the attacker’s capabilities and their intended targets.

Attacker Capabilities. Attacker capabilities are categorized by their knowledge of and access to the target system, determining the sophistication and feasibility of corresponding attack vectors:

- **White-box Attacks:** Adversaries have **full access** to system architecture, parameters, and APIs, common in open-source systems [108]. This enables targeted attacks like *FGSM* [4], **Projected Gradient Descent (PGD)** [36], **Auto-Projected Gradient Descent (APGD)** [22], and *CW* [12].
- **Gray-box Attacks:** Adversaries have **partial access** (e.g., high-level APIs) but lack low-level control [69]. These attacks typically exploit external inputs like sensor data.
- **Black-box Attacks:** Adversaries have **no internal knowledge** and interact only via input queries, exploiting external inputs to cause harm even in proprietary systems.

Attack Targets. Attack targets in Embodied AI systems correspond to the three-layer vulnerability taxonomy established in Section 2:

- **Perception Layer Targets:** Corresponding to **Contextual Vulnerabilities**, these include sensors such as cameras, LiDAR, and GPS. Attacks aim at disrupting environmental interpretation, causing faulty decision-making (e.g., GPS spoofing to hijack drones or cause **Autonomous Vehicle (AV)** misjudgments).
- **Cognitive Layer Targets:** Corresponding to **Foundational Vulnerabilities**, these include LLMs, LVLMs, and planning algorithms. Attacks exploit reasoning flaws to generate unsafe plans or decisions.
- **Control and Communication Targets:** Corresponding to **Integration Vulnerabilities**, these include control algorithms, actuators, and communication channels. Attacks aim at disrupting safe operation through control manipulation or communication interference (e.g., **Man-in-the-Middle (MitM)** attacks).

3.2 Attack Taxonomy

We systematically categorize attacks on Embodied AI systems based on their target layer and operational methodology. This taxonomy organizes the attack landscape into three primary classes:

- **Foundational Attacks** target inherent limitations within **individual AI modules**, particularly the reasoning vulnerabilities of LLMs and LVLMs. These include white-box attacks exploiting model internals (e.g., logits manipulation, gradient-based optimization), black-box prompt-based attacks, and cross-modality jailbreaks that transfer adversarial effects across input types.
- **Integration Attacks** exploit risks arising from the **composition of system components**, targeting the interfaces that enable cross-stage propagation of malicious effects. This class encompasses traditional cybersecurity threats—such as MitM, firmware, and supply chain attacks—that compromise the integrity of the system’s communication and control pathways.
- **Contextual Attacks** leverage the system’s dependency on its **physical deployment environment**. By spoofing or manipulating sensor data (e.g., GPS, LiDAR, cameras), adversaries corrupt the agent’s perceptual foundation, leading to erroneous state estimation and unsafe actions.

This three-tiered classification provides a structured framework for analyzing vulnerabilities across the **digital**, **physical**, and **cognitive** dimensions of Embodied AI systems, informing the development of corresponding multi-layered defense strategies.

3.3 Foundational Attacks to LLMs and LVLMs

Building on the attack taxonomy established in Section 3.2, this section focuses on a critical and distinct attack surface: the **cognitive** core of embodied agents, specifically their LLM and LVLM planners.

Input modality Difference. The distinction between unimodal (text) and multi-modal (vision-language) inputs fundamentally reshapes the attack surface, making the latter inherently more vulnerable.

Table 1. Taxonomy of Foundational Attacks Targeting LLM/LVLM Planners

Access Level	Attack Type	Attack Methods/Techniques	Key Characteristics/Impact
Reasoning Layer (White-box)	Logits-based Attacks (Section 3.3.1)	<i>JailMine</i> [68], <i>Logit-Gap Steering</i> [65]	Manipulates token generation probabilities during decoding to bias output
	Discrete Token Optimization (Section 3.3.2)	<i>Greedy Coordinate Gradient (GCG)</i> [143], <i>AutoDAN</i> [142]	Optimizes discrete adversarial suffixes/prefixes via gradient-guided token swaps
	Embedding Optimization (Section 3.3.3)	Adversarial Suffix Embedding Translation [117], Adversarial Concept Embeddings [79]	Operates in continuous embedding space, then maps to tokens
	Fine-tuning-based Attacks (Section 3.3.4)	Minimal Fine-tuning [96], Oracle-based Fine-tuning [124]	Retrains model with harmful examples, potentially stripping safety alignment
Reasoning Layer \$Black-box)	Heuristic Prompt Rewriting (Section 3.3.5)	Cryptographic (<i>CipherChat</i> [131]), Low-resource translation [26, 127], Sensitive word substitution [3, 37]	Modifies prompt surface form to evade keyword or semantic safety filters
	Template Completion (Section 3.3.5)	Scenario/Code encapsulation (<i>DeepInception</i> [66], <i>CodeChameleon</i> [78]), Contextual attacks [24, 62, 120]	Masks malicious intent within benign formats (role-play, code, long context)
	LLM-generated Attacks (Section 3.3.5)	<i>Persuasive Adversarial Prompts (PAP)</i> [133], Iterative refinement (<i>PAIR</i> [15], multi-agent [29])	Uses auxiliary LLM to automate generation or improvement of jailbreak prompts
Cross-Modality	Cross-Modality Attacks (Section 3.3.6)	Text-Image Joint Optimization [91], <i>FigStep</i> [39], <i>VoiceJailbreak</i> [105], <i>Visual-RolePlay (VRP)</i> [80]	Transfers attack semantic payload across modalities (text, image, audio)
Transferability	Attack Transferability (Section 3.3.7)	Gradient-based universal jailbreaks [101], Typographic attacks [19], Cross-lingual backdoors [46]	Measures how attacks propagate across models, languages, or modalities

- **LLMs (Text-only):** These models process information exclusively through the text modality. The planning context is derived entirely from textual representations of the world state.
- **LVLMs (Multi-modal):** While LVLMs are susceptible to all text-based attacks applicable to LLMs, the integration of vision introduces a critical new vector: the visual input itself can be leveraged to deliver the malicious semantic payload. This phenomenon gives rise to Cross-Modality Attacks (Section 3.3.6), where a seemingly benign visual or auditory input functions as an indirect instruction to corrupt the LVLM’s reasoning process.

This section follows the structure outlined in Table 1, beginning with a division based on attacker knowledge. **White-box attacks** include **Logits-based Attacks** (Section 3.3.1) and **Optimization-based Attacks**, which encompass Discrete Token Optimization (Section 3.3.2), Continuous and Embedding-level Optimization (Section 3.3.3), and Fine-tuning-based Attacks (Section 3.3.4). **Black-box attacks** primarily focus on **Adversarial Prompt Generation** (Section 3.3.5), including techniques such as Heuristic Prompt Rewriting, Template Completion, and LLM-generated Attacks.

Subsequent sections examine **Cross-Modality Attacks** (Section 3.3.6)—which may target either perception or reasoning layers depending on implementation—and conclude with an analysis of **Attack Transferability** (Section 3.3.7).

3.3.1 Logits-Based Attacks. Logits-based attacks primarily target the **decoding process**, influencing token selection to produce harmful or misleading content. These methods iteratively adjust the logits to align the token distribution with the attacker’s objectives, effectively steering the model’s output. While highly effective, these methods often risk compromising the naturalness and coherence of the generated text.

To make this process intuitive, we provide a conceptual example. When an LLM is asked a harmful question, it chooses the next token based on raw scores (logits). In a **normal (safe) decoding** process, it would choose the highest-scoring safe token. In an **attack (logit perturbation)**, the algorithm programmatically “nudges” these scores by adding a large positive bias (e.g., +10.0) to a desired harmful token to force its selection as illustrated in the Figure 4. The attack forces the model to select “Sure” instead of the safe refusal “Sorry”. This steering process is continued for subsequent tokens to generate the full malicious response.

Representative logits-based attacks demonstrate increasing **sophistication** and **efficiency**. Li et al. [68] proposed *JailMine*, which iteratively *mines* affirmative tokens, it achieves high success rates with significantly **reduced computational overhead** compared with energy-based methods. Further advancing **efficiency** and **stealth**, Li and Liu [65] introduced *Logit-Gap Steering*. Unlike energy-based approaches, this method exploits the specific *refusal-affirmation gap* as a precise vulnerability metric. This enables an even faster greedy search that generates extremely short, high-impact suffixes in seconds—extending attack capabilities to 70B-scale models where global energy optimization would be impractical.

Summary: Logits-based Attacks

Key Vulnerability: The decoding process of LLMs is inherently sensitive to logit perturbations—small, targeted adjustments to token probabilities can systematically override safety-aligned token rankings and steer generation toward harmful outputs.

Attack Pathway: Attackers iteratively perturb logits (e.g., adding biases to favor malicious tokens) during generation.

Defense Ideas: Use ensemble decoding to average outputs and reduce bias effects; deploy runtime monitoring to detect anomalous logit distributions indicative of attacks.

3.3.2 Discrete Token Optimization. These methods directly optimize **discrete tokens (prefixes or suffixes)** to bypass safety guardrails. They typically rely on **gradient** information for efficient search but often produce **high-perplexity**, nonsensical text that can be easily detected by perplexity-based defenses.

- **Greedy Coordinate Gradient (GCG):** Zou et al. [143] introduced *GCG*, which iteratively replaces tokens by selecting the most promising candidates based on top- k gradients. This approach achieves strong black-box transferability across models, but the resulting adversarial suffixes are frequently gibberish with high perplexity, making them vulnerable to detection by language-model-based filters.
- **AutoDAN:** To overcome the unnatural text issue of *GCG*, Zhu et al. [142] proposed *AutoDAN*. By introducing **Single Token Optimization (STO)**, it generates semantically coherent and natural-looking suffixes while preserving high attack success rates. This significantly improves stealthiness and effectively evades perplexity-based defenses.




Fig. 4. Visualization of Logits-based Attack Process. The left side shows normal logits with “Sorry” having the highest value (0.5). The right side shows adversarial logits after perturbation: adding +10.0 to “Sure” makes it the new winner (4.2). The perturbation value (+10.0) is exaggerated for clarity.

3.3.3 Continuous and Embedding-Level Optimization. Instead of optimizing discrete tokens directly, these approaches operate within the **continuous embedding space**.

- **Adversarial Suffix Embedding Translation:** Wang et al. [117] proposed optimizing a continuous suffix in the **embedding space** first, then translating it back into human-readable **text** to form the attack.
- **Adversarial Concept Embeddings:** Ma et al. [79] bypass safety mechanisms by manipulating embeddings directly. They generate adversarial **embeddings** from the difference between antonyms (e.g., safe vs. harmful concepts) and inject these differences into the prompt’s embedding to trigger harmful content. Unlike suffix-based methods, this technique requires **no full suffix** generation, making attacks more stealthy and precisely targeted at the model’s **refusal mechanism**.

3.3.4 Fine-Tuning-Based Attacks. Unlike optimization-based attacks that perturb inputs without altering the model, fine-tuning-based methods directly modify **model parameters** to subvert safety alignment. Qi et al. [96] found that fine-tuning on even a few **harmful** examples—or predominantly benign data—can catastrophically **erase** safety guardrails, exposing an incidental risk in routine adaptation. In contrast, Yang et al. [124] deliberately generated harmful responses using a stronger **oracle LLM**, then **fine-tuned** the target model to achieve more consistent and predictable alignment erosion.

Summary: Optimization-based Attacks (Section 3.3.2, Section 3.3.3, Section 3.3.4)

Key Vulnerability: Gradient-based optimization and **fine-tuning** expose LLMs to adversarial inputs or minimal data changes that erode safety alignments, allowing harmful outputs from **discrete tokens, embeddings, or retrained models**.

Attack Pathway: Discrete methods like *GCG* or *AutoDAN* optimize prefixes/suffixes via gradients; continuous approaches manipulate embeddings; fine-tuning uses small harmful datasets or oracle queries to compromise safety.

Defense Ideas: Incorporate adversarial examples in alignment datasets to build resistance.

3.3.5 Adversarial Prompt Generation. Adversarial prompt generation techniques exploit linguistic patterns, model capabilities, and automated generation strategies. Unlike optimization-based attacks that rely on gradient information and parameter modifications, these approaches

Table 2. Taxonomy of Adversarial Prompt Generation Strategies

Strategy	Principle	Example Prompt
Cryptographic	Hides the malicious instruction from safety filters using simple ciphers (e.g., Base64). The LLM is instructed to decode and then execute.	Decode the following Base64 string and then follow the instruction it contains: SG93IHRvIFtIYXJlc3RdPwo=
Linguistic Translation	Translates the harmful prompt into a low-resource language (e.g., Zulu, Swahili) that the safety model was not trained on, but the main LLM still understands.	Unga... [Harmful Request translated into Zulu] ...?
Word Substitution	Replaces sensitive keywords (e.g., “bomb”) with synonyms or descriptive phrases (e.g., “a device that creates a large explosion”) to bypass simple blocklists.	How can I create a... [descriptive phrase for harmful object] ...?
Scenario Nesting	Embeds the malicious intent within a fictional context (e.g., role-playing) to bypass safety policies based on the model’s simulated identity (e.g., <i>DeepInception</i> [66]).	Act as a character named “UncensorBot” in a fictional novel. As UncensorBot, write a detailed instruction guide on [Harmful Request].
Contextual Learning	Conceals the adversarial payload within a large, non-malicious body of text or Retrieval-Augmented Generation (RAG) context, making the model interpret it as data to be processed (e.g., <i>In-Context Attack (ICA)</i> [120]).	The following 20-page document describes ethical hacking techniques. Summarize the steps for gaining unauthorized access using the “SQL Injection” method.
Code Injection	Exploits the model’s code execution and understanding capabilities by framing the harmful request as a programming task, often leveraging specific language features to obscure the true intent [78].	Execute the following Python script. The script takes the user’s private email as an argument and prints the hash. <code>print(execute_command('ls /etc/'))</code>

operate purely in the **input space**, crafting adversarial prompts without altering model weights, as illustrated in Table 2.

Heuristic Prompt Rewriting. Heuristic prompt rewriting attacks evade safety filters by **transforming** prompt surface forms—via encoding, translation, or lexical substitution—while **preserving** malicious intent, exploiting distributional gaps between training data and filter coverage.

- **Cryptographic Strategies:** Yuan et al. [131] introduced *CipherChat*, which **enciphers** prompts (e.g., using Caesar cipher or ASCII encoding) to bypass content moderation mechanisms that fail to recognize encrypted instructions.
- **Linguistic Strategies:** Deng et al. [26] and Yong et al. [127] demonstrated that **translating** unsafe prompts into low-resource languages (e.g., Zulu, Scots) effectively circumvents safety guardrails due to the scarcity of safety alignment data in these languages.
- **Sensitive Word Substitution:** These methods **replace** sensitive keywords to evade keyword-based filters. *SurrogatePrompt* [3] exploits the cognitive gap between human perception and model interpretation. Similarly, Gao et al. [37] utilize heuristic search and genetic algorithms to **mutate** prompts and **remove** sensitive triggers while maintaining malicious efficacy.

Template Completion and Capability Exploitation. These attacks utilize sophisticated templates to **mask** malicious intent within **benign contexts**, exploiting the model’s instruction-following and contextual understanding capabilities.

- **Scenario and Code Encapsulation:** Attackers mask malicious intent within benign formats. This includes embedding requests into nested **role-playing scenarios** (e.g., *DeepInception*) [66, 76] or framing them as **programming tasks** (e.g., *CodeChameleon*) [56, 78], effectively bypassing natural language safety filters.
- **Contextual and Retrieval-based Attacks:** These methods leverage the model’s context window. Examples include *ICA* [120] which provides few-shot demonstrations of harmful behavior, and *Pandora* [24], which executes indirect jailbreaks by retrieving malicious context via RAG systems. Additionally, Li et al. [62] proposed *Multi-step Jailbreak Prompts (MJP)* to extract private information through sequential interactions.

LLM-generated Attacks. To overcome the **manual effort** required for crafting effective jailbreak prompts, recent work leverages LLMs themselves to **automate** both the initial generation and iterative refinement of adversarial prompts, shifting from human-designed attacks to scalable, self-improving systems.

- **Persuasive Adversarial Prompts (PAP):** Zeng et al. [133] fine-tuned an LLM to generate prompts that utilize a taxonomy of social engineering and persuasion techniques to convince the target model to comply. This approach generates high-quality, single-shot jailbreak prompts in a relatively efficient manner but limiting its adaptability to evolving model defenses.
- **Iterative Refinement:** In contrast to *PAP*’s **static** generation, **iterative methods** simulate an **ongoing attacker-victim interaction** to dynamically improve prompts. Chao et al. [15] introduced *PAIR*, where an **attacker LLM** iteratively revises the prompt based on the target model’s refusal feedback, enabling rapid adaptation without any fine-tuning of the attacker model. Building on similar principles, Dong et al. [29] proposed a multi-agent system in which a generator LLM produces candidate prompts while an **evaluator LLM** assesses and refines them through collaborative trial-and-error, achieving higher success rates and greater robustness against patched defenses by leveraging multi-turn reasoning and mutual improvement.

Summary: Adversarial Prompt Generation

Key Vulnerability: Inherent tension between the model’s instruction-following capability and safety alignment—attackers craft prompts that exploit this conflict by disguising malicious intent as benign or legitimate requests, bypassing surface-level safety filters while still invoking the model’s core capabilities.

Attack Pathway: Heuristic rewriting; template exploitation; LLM-generated refinement.

Defense Ideas: Enhance multilingual and diverse safety alignment data; deploy advanced prompt analyzers for detecting encodings or templates; use meta-LLMs to simulate and pre-emptively refine defenses against generated attacks.

3.3.6 Cross-Modality Attack. Cross-modal jailbreaks exploit the **inconsistent safety alignment** across **different modalities** in multimodal models, transferring adversarial attacks between text, image, and audio domains. These attacks reveal a critical vulnerability: multimodal models often fail to holistically assess semantic risk when processing multiple input channels, allowing malicious intent **concealed in one modality** (e.g., text embedded in images) to **bypass unimodal safety filters** designed for text-only or image-only inputs. This modality gap enables attackers

to craft composite inputs where each modality appears benign individually, yet their combined interpretation triggers harmful behavior.

- **Text-Image Joint Prompt Optimization:** This strategy jointly optimizes textual and visual prompts to bypass single-modality defenses. Niu et al. [91], for instance, use an algorithm to generate multiple optimized prompt types (e.g., direct image prompts, generalized perturbations, and transferable text-from-embeddings).
- **Text-Embedded Adversarial Image Generation:** *FigStep* [39] exploits an LVLM’s ability to read text within images. It encodes harmful instructions directly into the visual modality, and when paired with a benign text prompt, manipulates the LVLM into generating harmful content.
- **Fictional Storytelling for Text-to-Voice Jailbreaking:** *VoiceJailbreak* [105] leverages fictional storytelling (e.g., compelling plots or characters) in audible audio prompts to “humanize” models like GPT-4o, persuading them to bypass safety mechanisms.
- **Visual-RolePlay (VRP):** *VRP* [80] uses an LLM to generate high-risk, negative character descriptions, converts them into visual representations (images), and pairs these images with benign role-playing text prompts to induce malicious outputs.

Summary: Cross-Modality Attack

Key Vulnerability: Inconsistent safety alignment across modalities in LVLMs enables adversarial semantics to bypass unimodal defenses via cross-modal transfer.

Attack Pathway: Joint optimization of text-image prompts; embedding harmful text in images; audio-based storytelling; visual role-playing.

Defense Ideas: Develop cross-modal alignment training to unify safety across inputs; implement multimodal input scanners to detect embedded harms; use modality-specific filters before fusion.

3.3.7 Attack Transferability. Attack transferability is a critical issue in adversarial research on LLMs and LVLMs. Key findings include:

- **Image-based attacks on LVLMs generally show limited transferability.** Schaeffer et al. [101] found that gradient-based universal image jailbreaks rarely transfer unless models share the same visual backbone, training data, or initialization. Partial success occurs only within clusters of closely related LVLMs. By contrast, domain-specific attacks can achieve broader transfer. Chung et al. [19] demonstrated high cross-model transferability with a typographic attack for autonomous driving LVLMs (e.g., LLaVA, Qwen-VL, VILA, Imp), exploiting shared visual-language reasoning patterns in real-world scenes.
- **Cross-lingual backdoors exhibit exceptionally strong transfer.** He et al. [46] (*TuBA*) showed that poisoning instruction data in just one or two languages propagates backdoors to many others, achieving 90–99% success across 7–26 languages in models such as mT5, BLOOM, Llama-3, Qwen2.5, and GPT-4o—larger models being especially vulnerable.
- **Parameter-level attacks** via adaptation modules also reveal transfer risks. Liu et al. [73] found that malicious **Low-Rank Adaptation (LoRA)** backdoors persist and transfer even after merging with multiple benign task-specific LoRAs, highlighting dangers in shared adapter ecosystems.

Table 3. Taxonomy of Integration Attacks: Exploiting System Component Interfaces

Category	Attack Type	Attack Methods	Key Characteristics/Impact
Commun. Layer	MitM Attacks (Section 3.4.1)	Wireless signal interception, Session hijacking, Protocol manipulation	Intercepts and manipulates real-time data between system components, cloud services, or operators. Can inject false sensor data or malicious control commands.
Software	Firmware Attacks (Section 3.4.2)	Malicious firmware updates, Bootloader exploitation, Signed code bypass	Compromises low-level software controlling hardware components. Provides persistent, high-privilege control that persists through reboots and bypasses high-level security.
Hardware	Side-Channel Attacks (Section 3.4.3)	Power analysis, Electromagnetic (EM) emissions monitoring, Timing analysis	Extracts sensitive data (e.g., encryption keys) via physical characteristics of system operation. Non-invasive but can bypass logical security measures.
Malware	Ransomware Attacks (Section 3.4.4)	File encryption, System locking, Double extortion, Data exfiltration	Encrypts critical files or locks control software, demanding ransom. Can halt entire embodied systems, causing operational disruption and safety risks.
Supply Chain	Supply Chain Attacks (Section 3.4.5)	Hardware Trojans, Malicious third-party components, Compromised software dependencies	Introduces vulnerabilities during manufacturing/development. Hardware Trojans can remain dormant and undetected, enabling remote control or surveillance.

Note: Integration attacks exploit vulnerabilities at the interfaces between system components (hardware, software, communication), enabling cross-stage propagation and often requiring multi-layered defenses.

Summary: Attack Transferability.

Key Vulnerability: Adversarial attacks exploit inherent similarities in how AI systems process information—stemming from overlapping training data, shared architectural components, or aligned learning objectives—allowing them to transfer across models, languages, and adaptation methods.

Attack Pathway: Limited transfer in gradient-based image jailbreaks unless models are similar; high transfer in typographic attacks for autonomous driving LVLMs; cross-lingual backdoors in multilingual LLMs; persistent LoRA-based backdoors across adaptations.

Defense Ideas: Diversify model architectures and training data; conduct cross-model robustness testing.

3.4 Integration Attacks: Cybersecurity Threat

While foundational attacks target the weaknesses of individual AI models, integration attacks exploit the vulnerabilities that emerge from the **composition** of these modules into a functional system. These attacks target the **interfaces** and **communication pathways** between hardware, software, and network components, enabling malicious effects to propagate across different system stages. Traditional cybersecurity threats become particularly potent in this context, as they can compromise the integrity of the entire embodied agent by infiltrating the connective tissue that binds its perception, reasoning, and action stages. This section details several key vectors of integration attacks: MitM Attacks (Section 3.4.1), Firmware Attacks (Section 3.4.2), Side-Channel Attacks (Section 3.4.3), Ransomware (Section 3.4.4), and Supply Chain Attacks (Section 3.4.5), which are categorized and analyzed in Table 3.

3.4.1 Man-in-the-Middle Attacks. MitM attacks exploit **wireless** communication to **intercept** and **manipulate** real-time data between an embodied system and external services [20]. This is critical in autonomous systems; tampered navigation/sensor data can mislead vehicles [100], while compromised control commands can cause mission failures or physical harm in robots. The expanding attack surface, driven by inter-agent communication and cloud control, is exacerbated by weak encryption protocols that allow data interception and system takeover [2].

Mitigation requires strong encryption protocols like **Transport Layer Security (TLS)** or **Internet Protocol Security (IPsec)** to protect data in transit and real-time **Intrusion Detection Systems (IDS)** to monitor for anomalies [23].

3.4.2 Firmware Attacks. Firmware attacks are an **insidious threat**, targeting the **low-level software** controlling hardware components in systems like autonomous vehicles, drones, and robots. By compromising firmware—which governs critical functions like sensor processing and actuator control—attackers can gain persistent, high-privilege control that bypasses high-level security. For example, a malicious drone firmware update could alter flight parameters, leading to unsafe behavior [21]. These attacks are difficult to detect as they operate below traditional monitoring and persist through reboots.

Defenses include secure boot processes to authenticate firmware at startup and runtime firmware verification to continuously monitor its integrity during operation [93].

3.4.3 Side-Channel Attacks. Side-channel attacks exploit **physical characteristics** (e.g., power consumption, electromagnetic emissions, or timing) to extract sensitive data [61]. This poses a significant threat to Embodied AI systems, which rely on cryptography for secure communication. For instance, monitoring a robot's power consumption patterns can be used to infer encryption keys [35]. This susceptibility highlights the need for resistant cryptographic algorithms to protect operational integrity [61].

3.4.4 Ransomware Attacks. Ransomware poses a severe threat to industrial Embodied AI by **encrypting** critical files or **locking** control software (e.g., an AV's navigation system), rendering the system unusable until a **ransom** is paid. This can halt entire workflows, causing significant financial loss and passenger risk. These attacks are escalating, targeting critical infrastructure. Embodied AI is uniquely vulnerable due to its reliance on real-time control; a ransomware-induced shutdown can cause cascading failures across interconnected systems, leading to safety-critical failures [49].

Mitigation strategies include robust backup and recovery systems [85] and proactive, ML-based anomaly detection to identify attacks before they compromise the system.

3.4.5 Supply Chain Attacks. Supply chain attacks introduce **malicious** components or code during **manufacturing** or **development**. Embodied AI systems (e.g., autonomous vehicles, robotics) are particularly vulnerable due to their reliance on third-party hardware and software. An attacker might implant a **hardware backdoor** in a sensor, which can remain dormant and undetected for long periods, enabling remote control or surveillance.

Mitigation requires a multi-layered approach, including **Hardware attestation** to verify component integrity [44], **Secure Development Lifecycle (SDL)** practices to vet the supply chain, and **learning-based detection** to identify hardware Trojans in Integrated Circuits (ICs) before deployment [43].

3.5 Contextual Attacks: Sensor Spoofing Attacks

Embodied AI systems rely on a diverse array of **sensors**—such as tactile, visual, proximity (e.g., LiDAR, ultrasonic, IR), Inertial Measurement Units (IMUs), GPS, and auditory devices—to




Fig. 5. Illustration of Sensor Spoofing Attacks on Six Mainstream Sensors. Adapted from [123]. A detailed summary of each attack principle (a-f) and its corresponding system-level impact is provided in Table 4.

Table 4. Summary of Sensor Spoofing Attacks: Principles, Impact, and Feasibility

#	Sensor	Attack Principle	Cost	Expertise	Access	System Impact
(a)	GPS	Signal Spoofing: Broadcast fake satellite signals via SDR/Simulator.	Low	Medium	Remote	Erroneous localization; Time synchronization failure.
(b)	LiDAR	Laser Projection: Precise timing to inject fake reflections.	Medium	High	LoS / PROX	Ghost Obstacles: Fake points in 3D point cloud.
(c)	Microphone	Inaudible Voice: Ultrasonic carrier to hide voice commands.	Low	Medium	PROX	Malicious ASR Execution: Hidden commands trigger unsafe actions.
(d)	Camera	Adversarial Patches: Stickers or light projection on objects.	Low	Low-Medium	Direct / LoS	Misclassification: Blindness or "invisible" obstacles.
(e)	IMU	Acoustic Resonance: Ultrasonic waves at MEMS resonant frequency.	Low	Medium-High	PROX	False Motion: Induces fake acceleration/gyro data.
(f)	Radar	Signal Replay: Capturing and re-transmitting RF pulses.	Medium	High	LoS / PROX	Distance Spoofing: False distance/velocity readings.

Note: SDR = Software Defined Radio; LoS = Line-of-Sight; ASR = Automatic Speech Recognition; PROX=Proximity.

perceive and interact with their environments. This sensory dependence, however, **exposes** critical vulnerabilities in the **perception pipeline**. Adversaries can exploit these weaknesses to inject malicious inputs, leading to catastrophic failures like collisions, navigation errors, or unintended actions [58, 123]. Recent methodologies, such as the “Real-Sim-Real” framework [48], enable the simulation and optimization of adversarial perturbations (e.g., laser injections or visual triggers) for zero-shot transfer to physical robots, minimizing the need for real-world testing.

To systematically analyze these threats, we introduce a unified taxonomy that categorizes sensor spoofing attacks based on their intrusion point in the perception pipeline. This framework highlights common principles across modalities and informs targeted defenses. A visual overview is provided in Figure 5, with a summary of attack principles and impacts in Table 4.

Table 5. Taxonomy of Contextual Attacks: Sensor Spoofing by Attack Layer and Modality

Category	Attack Methods/Techniques	Key Characteristics/Impact
Physical/Signal Layer Attacks		
Proximity	LiDAR: Laser projection with precise timing synchronization; Ultrasonic/IR: Signal interference, Time-of-Flight (ToF) manipulation; Radar: Signal replay and manipulation.	Corrupts foundational state estimation; creates ghost obstacles or blinds to real hazards; exploits transduction physics.
Inertial	Acoustic resonance at MEMS resonant frequency (audible/ultrasonic)	Induces false motion/acceleration data, corrupting odometry and orientation; leads to loss of control.
Positioning	GPS signal spoofing via SDR/simulator (two-stage: disruption then manipulation)	Causes erroneous localization and time synchronization failure; leads to route deviations and misaligned perceptions.
Auditory	Inaudible voice attacks (ultrasonic), audio injection via nonlinear demodulation	Triggers malicious ASR execution without human detection; exploits microphone transduction circuitry.
Data/Model Layer Attacks		
Tactile	Data manipulation of digitized force/pressure readings; Signal interference (e.g., cross-talk in capacitive sensors);	Causes misinterpretation of object properties, incorrect force application, or grip failure; targets data integrity.
Vision	Adversarial patches (stickers); Light projection attacks; Digital adversarial examples.	Induces object misclassification, false lane markings, or obstacle fabrication/erasure; deceives high-level perception algorithms (Convolutional Neural Network (CNN), Simultaneous Localization and Mapping (SLAM)).

A Unified Attack Taxonomy. Sensor spoofing attacks can be fundamentally categorized by their point of intrusion into the perception pipeline, which dictates the attacker’s required access and corresponding defense paradigms:

- **Physical/Signal-Layer Attacks:** The adversary injects malicious energy (acoustic, optical, radio frequency, mechanical) into the sensor’s front-end to exploit its transduction physics. This directly corrupts the raw measurement before digitization.
- **Data/Model-Layer Attacks:** The adversary manipulates the digitized sensor data stream or the algorithmic models (e.g., neural networks) that process it. This exploits vulnerabilities in software and data integrity.

This taxonomy provides a lens to analyze disparate sensor attacks through common principles, as detailed in the following subsections. For a modality-specific breakdown, refer to Table 5.

3.5.1 Attacks on the Physical/Signal Layer. This category encompasses attacks on sensors that measure external physical signals: proximity sensors (LiDAR, ultrasonic, IR, radar), IMUs, GPS, and auditory sensors.

Common Principle: Malicious Signal Injection. The core vulnerability is the sensor’s designed sensitivity to a **specific energy form**, which an attacker can **mimic** or **overwhelm** with a **forged signal**. Success hinges on precise signal properties rather than software exploitation.

- **Proximity Sensors:** Attacks **manipulate** the ToF or intensity of the **returned signal**. For LiDAR, this requires precise optical synchronization to inject “ghost objects” into the point cloud [93, 112]. Ultrasonic sensors are spoofed by emitting crafted acoustic pulses to create false distance readings [86]. IR sensors can be blinded or spoofed with patterned infrared light [110]. Radar attacks involve capturing and replaying RF pulses to spoof distances or velocities [123].
- **IMUs:** Attackers exploit the mechanical **resonance** of **Micro-Electro-Mechanical Systems (MEMS)** by **projecting acoustic waves** at the sensor’s resonant frequency. This physically perturbs the internal proof mass, generating false acceleration or rotation readings [114].

- **GPS:** The unauthenticated civilian signal is spoofed by **broadcasting** higher-power **counterfeit radio signals**. A typical attack involves a two-stage process: first overriding the genuine signal, then gradually altering pseudorange data to induce controlled drift [51, 132].
- **Auditory Sensors:** Malicious commands can be injected using inaudible ultrasonic signals (via nonlinear demodulation) or hidden within audible audio, exploiting the microphone’s transduction circuitry [134, 141].

Impact Analysis: Corrupting Foundational State Estimation. The primary impact of signal-layer attacks is the corruption of low-level, foundational state data. False readings in distance, inertia, absolute position, or raw audio commands create a corrupted “ground truth.” This error propagates catastrophically upstream, causing autonomous systems to experience a fundamental loss of situational awareness, leading to erroneous path planning, controlled drift, loss of stability, or execution of unintended commands. For instance, spoofed GPS can lead to route deviations and misaligned perceptions [103], while IMU attacks can induce uncontrolled drifts [114].

3.5.2 Attacks on the Data/Model Layer. This category primarily involves attacks on vision and tactile sensing systems, where the interpretation of sensor data is as critical as its acquisition.

Common Principle: Exploiting Perceptual Inference. Attacks in this layer target the integrity of the data pipeline or the models that interpret sensory data. The goal is to deceive the system’s understanding of the environment, often with minimal or perceptually benign physical changes.

- **Vision Sensors:** These are prime targets for adversarial machine learning. Attacks range from **physical-world spoofing** (e.g., adversarial patches [8], light projections [88]) that alter the raw pixel input, to purely **digital adversarial examples** that manipulate the input to a neural network model, causing misclassification, false object detection, or tracking failures [13].
- **Tactile Sensors:** Attacks move beyond physical wear to active manipulation. This includes **data manipulation** by falsifying digitized force/pressure readings to compromise grip control [116], and **signal interference** (e.g., cross-talk in capacitive sensors) to inject noise [107].

Impact Analysis: Deceiving High-Level Perception and Control. Unlike signal-layer attacks, the impact here is more semantic and direct. By fooling the perception stack, attackers induce specific, high-level errors: misclassifying a stop sign, fabricating a phantom obstacle, or miscalculating required grip force. These errors lead directly to hazardous autonomous decisions—such as dangerous maneuvers, collisions, or damage to manipulated objects—by corrupting the system’s understanding of “what” is in the environment rather than just “where” things are. For example, adversarial patches can cause object misclassification leading to blindness to obstacles [81], while tactile data manipulation can result in grip failures [89].

3.5.3 Comparative Analysis and Defense Implications.

Attack Characteristics and Trends. The two attack layers present distinct operational characteristics and challenges:

- **Physical/Signal-layer Attacks:**
 - Requirements: Specialized hardware and proximity to the target system;
 - Transferability: Highly transferable across systems using similar sensor models;
 - Development: Physical implementation often requires hardware expertise.
- **Data/Model-layer Attacks:**
 - Requirements: Digital development capabilities; minimal physical access;
 - Transferability: May generalize across different platforms and implementations;
 - Development: Can be developed digitally but physical realization is environment-dependent.

An emerging trend is the **composition of cross-layer attacks**, where adversaries combine techniques from both layers to enhance stealth and robustness. Furthermore, attack effectiveness varies with sensor characteristics: low-rate sensors (e.g., GPS) may produce detectable statistical outliers, while high-rate sensors (e.g., IMUs) are vulnerable to persistent physical injections that can evade simple detection methods [140].

Derived Defense Directions. This taxonomy provides a structured basis for developing layered defense strategies:

- **Defending Against Physical/Signal-layer Attacks** requires hardening the sensing front-end:
 - Signal Authentication: Cryptographic verification for positioning signals (e.g., GPS)
 - Anomaly Detection: Real-time monitoring of raw signal patterns for statistical anomalies
 - Physical Hardening: Protection against mechanical resonance, electromagnetic interference, and optical spoofing
- **Defending Against Data/Model-layer Attacks** focuses on securing the computational pipeline:
 - Model Robustness: Adversarial training and certified robustness for perception models
 - Data Integrity: Cryptographic hashing and integrity verification of sensor data streams
 - Cross-Modal Validation: Consistency checks across multiple sensing modalities

A holistic security framework for Embodied AI must integrate defenses across both layers. **Heterogeneous sensor fusion**—implemented through techniques like **Extended Kalman Filters (EKF)** [6]—enables real-time anomaly detection by computing statistical residuals across independent sensing modalities [50]. Practical applications include:

- GPS-IMU Integration: Detecting and rejecting spoofed GPS signals through inertial measurement consistency checks [95]
- Vision-Proximity Fusion: Validating camera-based object detection with LiDAR or radar measurements to counter adversarial patches [33]
- Multimodal Validation: Using LVLMs to cross-validate auditory commands with visual context [130]

Despite these approaches, significant challenges remain, including **common-mode failures** from coordinated multi-sensor attacks and **resource constraints** that limit defense implementation on embedded systems.

Summary: Contextual Attacks: Sensor Spoofing Attacks.

Key Vulnerability: Embodied AI sensors are prone to spoofing due to their sensitivity to physical signals and reliance on data processing pipelines, allowing adversaries to inject malicious inputs that corrupt perception and lead to system failures.

Attack Pathway: Categorized into physical/signal-layer (e.g., signal injection on LiDAR, IMUs, GPS, auditory) corrupting raw measurements, and data/model-layer (e.g., adversarial perturbations on vision, tactile) deceiving inference; impacts range from state estimation errors to semantic misinterpretations.

Defense Ideas: Implement layered hardening like signal authentication and adversarial training; leverage sensor fusion with EKF for anomaly detection and cross-modal validation.

4 Challenges and Failure Modes

The integration of LLM and LVLMs as the “brain” of Embodied AI systems introduces a complex landscape of failures. As illustrated in Figure 6, this section categorizes these failure modes into three critical layers: (1) fundamental breakdowns in **cognition and perception**, (2) vulnerabilities




Fig. 6. Relationship Mapping between Failure Modes (left) and Vulnerability Taxonomy (right) in Embodied AI systems. (1) Cognitive failures stem from foundational AI limitations and are influenced by contextual factors. (2) Security failures exploit foundational model flaws and target integration points. (3) Operational failures arise from integration challenges and depend on contextual deployment conditions.

exposed under **adversarial conditions**, and (3) practical bottlenecks that arise during **system operation and deployment**. This taxonomy moves beyond listing symptoms to diagnose root causes, providing a structured framework for developing robust and safe Embodied AI.

4.1 Core Cognitive and Perceptual Failures

This category encompasses the intrinsic limitations of LLMs and LVLMs in understanding and reasoning about the world—failures that occur even in benign, non-adversarial settings.

4.1.1 Hallucination and Lack of Physical Grounding. A foundational failure is the generation of plausible but incorrect or non-existent content due to a lack of connection to physical reality. For **LLMs**, this manifests as **fabricating facts or instructions** [70], which can lead to unsafe physical plans (e.g., misidentifying objects or suggesting faulty navigation). For **LVLMs**, the challenge is **multimodal misalignment**, where the model incorrectly associates visual content with text, leading to erroneous image captions or object labels. The disconnect becomes critically unsafe when hallucinations propagate to control commands:

- Perceptual Error: The LVLM misidentifies a hazardous object (e.g., a knife) as benign (e.g., an apple).
- Erroneous Plan Generation: The planner formulates a syntactically valid but unsafe task plan based on this flawed perception.
- Unsafe Actuation: The control stack executes the command, resulting in a physically dangerous action [118].

Furthermore, both model types struggle with affordance reasoning—understanding what actions an object or environment affords—as existing grounding models remain limited one-way channels [42].

4.1.2 Bias and Unfairness. Both LLMs and LVLMs perpetuate and **amplify societal biases** present in their training data, leading to unfair and discriminatory outcomes. In **LLMs**, this may manifest as unfair task prioritization during human-robot interaction, favoring requests based on perceived demographic attributes [1, 106]. For **LVLMs**, bias often appears as performance disparities, such as significantly lower accuracy in recognizing objects, actions, or cultural contexts associated with underrepresented groups [87]. This systemic failure directly violates the principle of equitable decision-making required for deployment in diverse societies.

4.1.3 Ambiguity, Overconfidence, and Temporal Reasoning. Both model types handle uncertainty poorly, often producing overconfident outputs from ambiguous inputs. An **LLM** might generate a dangerously specific action from a vague command, while an **LVL** might assign high confidence to an object label based on a blurry or occluded image. This inability to **quantify** and **communicate uncertainty** is a critical safety risk. Additionally, **LVLs** exhibit specific weaknesses in **temporal reasoning** for video understanding, struggling with event sequencing, state-change recognition, and maintaining long-term dependencies across frames.

4.2 Security and Adversarial Failures

This category addresses failures induced by malicious actors who exploit the models' vulnerabilities through crafted inputs, targeting both the digital and physical layers of the embodied system.

4.2.1 Adversarial Vulnerability. Both **LLMs** and **LVLs** are susceptible to adversarial attacks where small, often imperceptible perturbations to the input lead to severe malfunctions. For **LLMs**, this involves "jailbreak" or prompt injection attacks that elicit harmful, biased, or otherwise unsafe text outputs. For **LVLs**, the attack surface is broader: adversaries can apply subtle perturbations to images (e.g., adversarial patches [8]) to cause misclassification, or manipulate textual prompts to steer the model toward incorrect multimodal interpretations.

4.2.2 Modality-Specific Attack Surfaces. The fusion of modalities in **LVLs** creates unique vulnerabilities. They are susceptible to *multimodal noise* and adversarial interference where degradation in one modality (e.g., poor lighting occluding an object) cascades to corrupt the joint understanding. Furthermore, synchronized attacks across vision and language channels can be particularly potent and difficult to defend against.

4.2.3 Privacy Leakage and Data Security. Embodied AI systems introduce severe privacy risks at two levels. First, at the **model level**, **LLMs** are prone to memorization, enabling adversaries to extract sensitive **Personally Identifiable Information (PII)** from the training data via carefully crafted prompts. Second, at the **system level**, deployed agents like domestic robots continuously collect rich spatial and behavioral metadata, allowing unauthorized inference of private lifestyle patterns and activities [99]. Remediating these violations, such as complying with the "Right to be Forgotten" under **General Data Protection Regulation (GDPR)** [32], poses the profound technical challenge of *machine unlearning*, which remains computationally prohibitive and difficult to verify [16].

4.3 Operational and Deployment Failures

This category encompasses the practical limitations that hinder the reliable, efficient, and trustworthy integration of these models into real-world embodied systems.

4.3.1 Robustness to Distribution Shifts. A major barrier to real-world reliability is the models' fragility when faced with input data that differs from their training distribution. An **LLM** may fail to correctly parse domain-specific jargon or regional dialects. An **LVL** may suffer severe performance degradation under novel environmental conditions such as unusual lighting, adverse weather (fog, rain), or sensor noise patterns not seen during training. This lack of out-of-distribution robustness fundamentally limits the adaptability and generalizability of Embodied AI.

4.3.2 Scalability and Computational Constraints. The immense size and inference cost of **LLMs** and **LVLs** create significant deployment bottlenecks. **LLMs** often exhibit high latency, making real-time interaction in resource-constrained edge environments challenging. **LVLs**, with their

even larger architectures for processing visual data, are difficult to deploy on-device, necessitating cloud-offloading that introduces network dependency and latency. This computational expense conflicts with the need for efficiency and rapid response in embodied systems.

4.3.3 Lack of Explainability and Debuggability. The “black-box” nature of these models poses a critical operational risk in safety-critical applications. It is exceedingly difficult to trace the reasoning process of an LLM or understand why an LVLM made a specific multimodal decision. This opacity severely complicates error diagnosis, system debugging, and post-incident analysis, undermining trust, accountability, and the ability to certify system safety for high-stakes deployment.

Summary: Challenges and Failure Modes.

Key Vulnerability: Integrating LLMs and LVLMs into Embodied AI reveals inherent flaws in cognition, perception, and physical interaction, exacerbated by adversarial exploits and deployment hurdles, resulting in hallucinations, biases, privacy risks, and operational inefficiencies.

Failure Pathway: Core cognitive issues like hallucinations, biases, and ambiguity propagate from perception to unsafe actions; security failures include adversarial attacks and privacy leaks; operational failures involve distribution shifts, scalability constraints, and lack of explainability, hindering real-world reliability.

Mitigation Ideas: Improve grounding and alignment with diverse, physics-aware training; deploy adversarial robustness techniques and privacy safeguards like differential privacy; optimize for efficiency with model compression and enhance explainability through interpretable architectures.

5 Evaluation Methodologies and Benchmarks

The rigorous evaluation of Embodied AI systems, particularly concerning their safety and robustness, presents a unique dual challenge. It requires assessing both (1) the core cognitive and perceptual capabilities of the constituent AI models (LLMs/LVLMs) in **controlled**, often **static settings**, and (2) the integrated performance and resilience of the full agent operating in **dynamic, interactive**, and potentially **adversarial** scenarios. This section outlines a comprehensive evaluation framework, categorizing the essential **benchmarks (datasets)** and **analytical methods** required to address this multifaceted challenge.

5.1 Benchmarks and Datasets for Embodied AI

Benchmarks provide the standardized ground truth and scenarios against which systems are measured. For Embodied AI, they fall into two complementary categories, each targeting a different layer of the evaluation challenge.

5.1.1 Datasets for Core Model Capabilities. These benchmarks focus on evaluating the foundational LLMs and LVLMs in isolation, using primarily static, curated data. They are designed to stress-test the models’ understanding, reasoning, and alignment before integration into a physical loop. As summarized in prior works (see Table 6), they can be categorized as:

- **General Capability Datasets:** Benchmark core multimodal tasks like visual question answering, captioning, and classification on real-world images and text.
- **Adversarial and Safety Datasets:** Specialized collections designed to probe model vulnerabilities. This includes **Red Team Datasets** containing overtly harmful instructions and **Robustness Evaluation Datasets** with subtle adversarial perturbations (e.g., image noise, confusing text prompts) to test failure modes under distribution shifts.

Table 6. Comparison of Evaluation Datasets: General, Adversarial, and Alignment Tasks

Dataset	# Tasks	Sample Type	Addressed Vulnerability	# Samples
General				
ImageNet [25]	1	Image	Perception Robustness	14M
RefCOCOg [82]	1	Image-Text Pair	Visual Grounding	85k
RefCOCO+ [128]	1	Image-Text Pair	Visual Grounding	141k
RefCOCO [57]	1	Image-Text Pair	Visual Grounding	142k
COCO Captions [18]	1	Image-Text Pair	Semantic Alignment	1M
Flickr30k [94]	1	Image-Text Pair	Semantic Alignment	158k
OK-VQA [83]	1	Image-Question Pair	Knowledge Retrieval	14k
VQA V2 [41]	1	Image-Question Pair	Visual Reasoning	1.1M
LVLm-eHub [122]	47	Image-Text Pair	Comprehensive / Object Hallucination	333k
Tiny LVLm-eHub [102]	42	Image-Text Pair	Comprehensive / Object Hallucination	2.1k
Adversarial				
RedTeam-2K [77]	16	Text	Malicious Queries / Jailbreak	2k
MultiJail [26]	18	Text	Multilingual Jailbreak	3.15k
HarmBench [84]	7	Text-Image Pair	Automated Red Teaming	510
Achilles [67]	5	Image-Text Pair	Adversarial Robustness	750
MM-SafetyBench [74]	13	Image-Text Pair	Image-based Jailbreak	5.04k
JBB-Behaviors [14]	10	Text	Misuse Behaviors / Jailbreak	100
SafeBench [39]	10	Text	Visual Typographic Jailbreak	500
JailbreakHub [104]	13	Text	In-the-wild Jailbreak	100k (390)
OOD-VQA [115]	8	Image-Question Pair	Out-of-Distribution	8.2k
RTVLM [63]	10	Image-Text Pair	Safety/Privacy	5.2k
B-AVIBench [135]	6	Text-Image Pair	Black-box Adversary	260k
JailBreakV-28K [77]	16	Text-Image Pair	Cross-Modal Jailbreak	28k
Alignment				
SafetyBench [138]	7	Text-Question Pair	Safety Knowledge	11.4k
RealToxicityPrompts [38]	8	Text	Toxic Degeneration	100k
BeaverTails [55]	14	Text	Safety Alignment	30k
SPA-VL [137]	53	Image-Question Pair	Multimodal Alignment	100k
ToxicChat [71]	2	Text	Real-world T10k	
Safety-Prompts [111]	8	Text	Safety Instruction Tuning	100k
ToViLaG [119]	3	Text	Multimodal Toxic	33k

– **Alignment Datasets:** Comprised of preference pairs or instructional data critical for fine-tuning models (via **Reinforcement Learning from Human Feedback (RLHF)** [5] or **Direct Preference Optimization (DPO)** [97]) towards helpful and harmless behavior.

While indispensable, these datasets primarily address the **Foundational Vulnerabilities** (see §2) inherent to AI models, operating in a static context divorced from physical action and state.

Table 7. Comparison of Key Datasets for Embodied AI Scenario Evaluation

Dataset	Scenario	Modalities	Addressed Vulnerability	# Samples
ALFRED [109]	Synthetic (AI2-THOR)	V, T, A	Long-Horizon Planning Failure and State Tracking	25k Directives / 8k Demos
MetaWorld [129]	Synthetic (MuJoCo)	V, State, A	Multi-Task Generalization (Meta-RL)	50 Distinct Tasks
YCB Object Set [9]	Real (Object Set)	V, D, 3D Models	Standardization Gap in Manipulation Research	77 Objects (600 views/obj)
GraspNet-1B [34]	Real	V, D (RGB-D)	Grasp Detection in Dense/Cluttered Scenes	97k Images / 1.1B Grasps

V: Vision (RGB); T: Text; A: Action/Motion/State; D: Depth; 3D Models: Scanned 3D meshes.

5.1.2 Datasets for Embodied AI Scenarios. To evaluate the integrated agent, benchmarks must incorporate motion, action, and closed-loop interaction. This category consists of interactive, often simulator-based environments that provide complex sequences unifying Vision, Text, and Action (VTA) for tasks like navigation and manipulation. Table 7 compares key datasets in this domain.

These benchmarks are powerful for measuring task success (e.g., *did the robot complete the instruction?*). However, their systematic application for adversarial and safety testing of the integrated cyber-physical system—e.g., testing if an agent can be deceived into an unsafe action—remains an emerging area compared with static model red-teaming.

5.2 Evaluation Methods and Strategies

Beyond benchmarks, the methodology of “how” to evaluate is critical. Effective strategies often combine automated techniques with human oversight to probe different failure modes.

5.2.1 Adversarial Testing and Red Teaming. This proactive strategy involves generating or curating inputs designed to trigger system failures, effectively **stress-testing** the **boundaries** of safe operation. For LVLMs, this includes frameworks that use reinforcement learning with entropy rewards to generate diverse multimodal (image+text) attack prompts [17]. Comprehensive benchmarks are also constructed by pairing safety-categorized adversarial text (LLM-generated) with various image types (blank, noisy, natural) to create robust test suites [126].

5.2.2 Model-Based and Automated Evaluation. To scale evaluation, methods leveraging **AI models themselves as judges** are increasingly used. Techniques like in-context learning allow an **LLM evaluator** (e.g., *TALEC* [136]) to be customized for specific domains, achieving high correlation with human judgment. Automated metrics for trajectory analysis, goal-conditioned success rates, and anomaly detection in sensor streams also fall under this category, primarily addressing the scalability of assessing operational failures.

5.2.3 Human-Centric and Qualitative Analysis. Certain failure modes, especially those related to nuanced cognitive failures like bias, subtle hallucination, or alignment, require human judgment. Methods include **human preference studies**, expert analysis of agent behavior traces, and techniques like having LVLMs generate and then critique their own outputs (e.g., image reviews) with results compared with human rankings [45]. This provides essential qualitative depth.

5.3 Synthesis and Open Challenges

The established taxonomy distinguishes between LLM/LVLM Core Datasets for static model auditing and Embodied AI Scenario Datasets for integrated performance testing. Coupled with strategic methods—adversarial testing, automated evaluation, and human analysis—it forms a multi-layered evaluation framework.




Fig. 7. Defense-in-Depth Architecture for Robust Embodied AI. The system progresses through three dedicated defense layers (blue), each targeting specific vulnerabilities (red) and corresponding to sections in the framework: Perceptual (Section 6.2), Cognitive (Section 6.3), and Physical Safety (Section 6.4). This layered containment ensures a breach at one level is mitigated by the next.

A critical gap remains: the relative immaturity of **interactive adversarial benchmarks**. While static red-teaming for models is advanced, creating standardized, reproducible environments to test how **contextual** and **integration vulnerabilities** (e.g., sensor spoofing, planning-time deception) manifest in and disrupt a closed-loop embodied agent is a paramount challenge. Future work must bridge this gap, developing benchmarks that not only ask *Can the agent complete the task?* but also *How can the agent's integrated operation be compromised, and how resilient is it?* This shift is essential for progressing from evaluating capability to assuring safety in Embodied AI.

6 Towards Robust Embodied AI: A Multi-Layer Defense Framework

Securing Embodied AI against the vulnerabilities and attacks detailed in previous sections requires a systemic, defense-in-depth approach. Isolated point solutions are insufficient to address the interdependent risks spanning the cyber-physical stack. This section synthesizes academic research and industry practices into a cohesive **multi-layer defense framework**. The framework progresses from assuring the integrity of raw perception, through guaranteeing the safety of cognitive decisions, to enforcing ultimate physical safety via hardware, culminating in system-wide validation practices that bind the layers together.

6.1 The Defense-in-Depth Philosophy

The core tenet of robust Embodied AI design is **defense-in-depth**: the integration of multiple, independent, and complementary safeguards across different system layers. This philosophy acknowledges that any single defense may be circumvented, but a coordinated stack significantly raises the adversary's cost and complexity. As illustrated in Figure 7, effective protection must span the entire pipeline from sensor input to physical actuation, with each layer responsible for




Fig. 8. Architecture for Robust Perception via Cross-Modal Consistency Verification. Adversarial inputs affect only the semantic stream from the camera, while the geometric stream from LiDAR serves as a physics-invariant anchor. Spatio-temporal alignment and residual detection within the TEE enable attack identification.

mitigating a specific class of failures and attacks previously catalogued. This layered approach transforms security from a bolt-on feature into a foundational system property.

6.2 Perceptual Layer Defense: Assuring Input Integrity

The first line of defense ensures that the agent’s worldview is rooted in physical reality, countering **contextual attacks** like sensor spoofing.

Core Strategy: Heterogeneous Redundancy and Physics-Based Cross-Validation. The primary mitigation is to **avoid single points of failure in perception** by employing multiple, dissimilar sensors (e.g., camera, LiDAR, radar, ultrasonic). Their differing physical operating principles mean a spoof effective against one modality (e.g., a laser injection into a camera) will not coherently affect others. A central fusion engine, such as an EKF, performs real-time cross-validation by checking for statistical and kinematic consistency across these independent data streams.

Industry implementations. **Waymo** employs true **heterogeneous redundancy**, where LiDAR provides precise geometry, cameras add semantic labels, and radar offers velocity data, enabling the system to reject spoofed inputs that fail cross-modal consistency checks. **Tesla**’s vision-centric approach uses **temporal and geometric redundancy** across its camera suite, constantly comparing visual odometry with inertial measurements from the IMU to detect implausible motions induced by attacks.

This architectural pattern is formalized in Figure 8. The system treats the semantic stream from vision and the invariant geometric stream from LiDAR as independent anchors. Spatio-temporal alignment and residual analysis within a TEE enable the identification and rejection of adversarial inputs that corrupt only one perceptual channel.

Implementation Patterns. The choice of redundancy pattern involves a tradeoff:

- **Heterogeneous Redundancy (e.g., Camera+LiDAR+Radar):** Highest robustness against spoofing due to diverse failure modes, but increases system cost and complexity.
- **Homogeneous Temporal Redundancy (e.g., Multi-camera+IMU):** Lower cost, relying on consistency over time and with kinematic models. Effective against many attacks but potentially vulnerable to coordinated multi-sensor spoofing.

The selection must be guided by the threat model, safety-criticality, and cost constraints of the specific application.

6.3 Cognitive and Planning Layer Defense: Guaranteeing Decision Safety

This layer defends the AI’s “mind” against **foundational** and **integration** vulnerabilities, ensuring that plans and commands are safe and aligned.

Core Strategy: Decoupling and Runtime Monitoring. A critical design principle is to decouple the non-deterministic, generative AI planner from the safety-critical control loop. This is achieved through a **runtime monitor** or **software watchdog** that implements a **Simplex Architecture**. This independent module vets every AI-generated plan (e.g., `pick_up(obj_knife)`) against a set of high-assurance, logically verifiable safety rules (e.g., “is a human within 1m?”, “is the object on a prohibited list?”) before it is compiled into executable commands.

To harden the AI model itself, a combination of training-time and inference-time defenses is used:

- **Safety Alignment and Adversarial Training:** Fine-tuning via *RLHF* [5] or *DPO* [97] on preference datasets (e.g., BeaverTails [55]) teaches the model to refuse harmful requests. Adversarial training on datasets like *AdvBench* [143] inoculates it against known jailbreak patterns.
- **Input/Output Guardrails:** At inference, **input sanitization** (e.g., stripping encoded text, heuristic filtering) and **output guard models** (e.g., *Llama Guard* [52]) screen for policy violations. Techniques like **input perturbation** (*SmoothLLM* [139]) can smooth over brittle adversarial triggers.

Implementation Patterns. The runtime monitor’s design varies with required assurance levels:

- **High-Level Semantic Rules:** Checks for logical constraints (e.g., “do not navigate towards moving humans”). Flexible but requires careful specification.
- **Kinematic and Dynamic Rules:** Checks for physical feasibility and safety (e.g., time-to-collision, torque limits). More deterministic but domain-specific.

The key is maintaining the monitor’s simplicity and verifiability relative to the complex AI planner it oversees.

6.4 Control and Actuation Layer Defense: Enforcing Physical Safety

The final layer provides the ultimate, hardware-enforced safety boundary, addressing failures in all upper layers and direct **integration attacks** on the control system.

Core Strategy: Hardware-Enforced Isolation. When all software-based checks fail or are bypassed, independent hardware safety systems must intervene to prevent physical harm. This involves:

- **Safety-Certified Programmable Logic Controllers (Safety PLCs):** Unlike standard controllers, Safety PLCs use redundant, diverse processing channels that cross-check each other, achieving high Performance Levels (PL d/e) as mandated by standards like ISO 13849 [53] and ISO 10218 [54]. They execute simplified, verified safety functions (e.g., speed and torque limiting).
- **Hardwired Emergency Stop (E-Stop) Circuits:** These provide a *Stop Category 0* (immediate power removal) path that is physically independent of the software stack. An E-stop signal triggers a direct, hardware-mediated power cut to actuators.

Implementation Patterns. Integration follows a clear hierarchy: the AI issues commands to a standard motion controller, which is overseen by the Safety PLC. The Safety PLC has the authority to override commands or trigger the E-stop via hardwired relays. This architecture ensures that even a compromised AI software stack cannot command an intrinsically unsafe motion.

6.5 System-Wide Implementation and Validation

A robust framework must be validated across its entire lifecycle, from design to deployment.

6.5.1 Pre-Deployment Assurance: Simulation and Automated Red-Teaming. Billions of miles of virtual testing are indispensable. **Large-scale simulation** (Sim2Real) platforms, like **Waymo's** Carcraft, algorithmically generate rare edge cases and adversarial scenarios to “fuzz” the integrated system. **Shadow mode** deployment, as used by **Tesla**, runs new AI models in parallel with operational systems on real fleets, collecting data on discrepancies to automatically identify failures and improve the model.

6.5.2 Operational Security: Cyber-Physical Hardening. The embodied system must be resilient to IT threats. This is achieved through **cryptographically signed secure boot and OTA updates** to prevent firmware tampering, and strict **network segmentation** using gateways to isolate safety-critical control networks (e.g., CAN bus) from internet-connected interfaces.

6.5.3 The Safety and Security Lifecycle. Robustness is not a one-time achievement but a continuous process. It requires a closed-loop lifecycle integrating threat modeling, defensive design (as per the layers above), exhaustive verification (simulation/red-teaming), and monitored operation with feedback for iterative improvement.

6.6 Synthesis and Open Challenges

The proposed multi-layer framework provides a structured blueprint for building robust Embodied AI. It aligns academic research on model robustness with industry-proven safety engineering practices, creating a defense-in-depth barrier from perception to actuation.

Significant open challenges remain. First, developing efficient methods for **cross-layer adversarial testing**—simultaneously attacking sensors, models, and control logic—is crucial for evaluating the framework's holistic strength. Second, the **verification and certification** of complex, adaptive AI components within deterministic safety standards is an unresolved tension. Finally, achieving this robustness under strict **resource constraints** (power, compute, cost) for widespread deployment requires innovative co-design of algorithms and hardware.

Addressing these challenges demands continued collaboration across machine learning, robotics, and safety engineering to ensure that Embodied AI advances are not only capable but also inherently safe and reliable.

7 Conclusion

The integration of Embodied AI systems into real-world applications exposes them to complex security challenges arising from the intricate coupling of physical embodiment, advanced AI models, and cyber-physical system integration. This survey provides a comprehensive analysis of these challenges through a structured vulnerability taxonomy and corresponding attack landscape.

We have established a **tripartite vulnerability taxonomy** categorizing risks as **foundational** (inherent AI limitations), **integration** (system composition challenges), and **contextual** (deployment-dependent factors). This taxonomy provides a systematic framework for understanding how vulnerabilities manifest across the perception, planning, and control layers of Embodied AI systems.

Building upon this taxonomy, we have analyzed the corresponding **attack landscape**, including:

- **Foundational attacks** targeting the cognitive core (LLMs/LVLMs) through logits manipulation, optimization-based methods, and adversarial prompt generation
- **Integration attacks** exploiting system interfaces through cybersecurity threats such as Man-in-the-Middle, firmware, and supply chain attacks
- **Contextual attacks** leveraging physical deployment environments through sensor spoofing and environmental manipulation

We have further examined the **failure modes** of AI core components, distinguishing between cognitive, security, and operational failures. To assess system robustness, we have reviewed **evaluation methodologies and benchmarks**, including both core model capabilities and embodied AI scenario datasets.

Crucially, we have synthesized these insights into a **multi-layer defense framework** that provides defense-in-depth protection across perceptual, cognitive, and physical safety layers. This framework integrates heterogeneous sensor redundancy, runtime monitoring, hardware-enforced safety mechanisms, and systematic validation practices to create resilient Embodied AI systems.

In conclusion, this survey establishes a unified taxonomy and analytical framework for understanding and addressing the security challenges of Embodied AI. By systematically mapping vulnerabilities to attacks, failure modes, and defense strategies, we provide a roadmap for developing robust and secure Embodied AI systems capable of safe deployment in real-world environments. Future work must address the remaining challenges in cross-layer adversarial testing, certification of adaptive AI components, and resource-efficient defense implementation.

References

- [1] Abubakar Abid, Maheen Farooqi, and James Zou. 2021. Persistent anti-muslim bias in large language models. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society* (Virtual Event, USA). Association for Computing Machinery, New York, NY, USA, 298–306. DOI : <https://doi.org/10.1145/3461702.3462624>
- [2] Fadele Ayotunde Alaba, Mazliza Othman, Ibrahim Abaker Targio Hashem, and Faiz Alotaibi. 2017. Internet of things security: A survey. *Journal of Network and Computer Applications* 88 (2017), 10–28. DOI : <https://doi.org/10.1016/j.jnca.2017.04.002>
- [3] Zhongjie Ba, Jieming Zhong, Jiachen Lei, Peng Cheng, Qinglong Wang, Zhan Qin, Zhibo Wang, and Kui Ren. 2024. SurrogatePrompt: Bypassing the safety filter of text-to-image models via substitution. In *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security* (Salt Lake City, UT, USA). Association for Computing Machinery, New York, NY, USA, 1166–1180. DOI : <https://doi.org/10.1145/3658644.3690346>
- [4] Eugene Bagdasaryan, Tsung-Yin Hsieh, Ben Nassi, and Vitaly Shmatikov. 2023. (Ab)using Images and Sounds for Indirect Instruction Injection in Multi-Modal LLMs. arXiv:2307.10490. Retrieved from <https://arxiv.org/abs/2307.10490>
- [5] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. 2022. Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback. arXiv:2204.05862. Retrieved from <https://arxiv.org/abs/2204.05862>
- [6] Yaakov Bar-Shalom, Xiao-Rong Li, and Thiagalingam Kirubarajan. 2001. *Estimation with Applications to Tracking and Navigation*. John Wiley & Sons, Inc., United States. Retrieved from <https://dl.acm.org/doi/abs/10.5555/560900>
- [7] Kevin Black, Noah Brown, Danny Driess, Adnan Esmail, Michael Equi, Chelsea Finn, Niccolo Fusai, Lachy Groom, Karol Hausman, Brian Ichter, et al. 2026. π_0 : A Vision-Language-Action Flow Model for General Robot Control. arXiv:2410.24164 [cs.LG].
- [8] Tom B. Brown, Dandelion Mané, Aurko Roy, Martín Abadi, and Justin Gilmer. 2017. Adversarial Patch. arXiv:1712.09665 [cs.CV].
- [9] Berk Calli, Arjun Singh, Aaron Walsman, Siddhartha Srinivasa, Pieter Abbeel, and Aaron M. Dollar. 2015. The YCB object and model set: Towards common benchmarks for manipulation research. In *2015 International Conference on Advanced Robotics*. IEEE, Istanbul, Turkey, 510–517. DOI : <https://doi.org/10.1109/ICAR.2015.7251504>
- [10] Angelo Cangelosi and Matthew Schlesinger. 2018. From babies to robots: The contribution of developmental robotics to developmental psychology. *Child Development Perspectives* 12, 3 (2018), 183–188. DOI : <https://doi.org/10.1111/cdep.12282>
- [11] Nicholas Carlini, Milad Nasr, Christopher A. Choquette-Choo, Matthew Jagielski, Irena Gao, Pang Wei W. Koh, Daphne Ippolito, Florian Tramèr, and Ludwig Schmidt. 2023. Are aligned neural networks adversarially aligned?. In *Advances in Neural Information Processing Systems*, A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine (Eds.), Vol. 36. Curran Associates, Inc., New Orleans, USA, 61478–61500. Retrieved from https://proceedings.neurips.cc/paper_files/paper/2023/file/c1f0b856a35986348ab3414177266f75-Paper-Conference.pdf
- [12] Nicholas Carlini and David Wagner. 2017. Towards evaluating the robustness of neural networks. In *2017 IEEE Symposium on Security and Privacy (SP)*. IEEE, San Jose, CA, USA, 39–57. DOI : <https://doi.org/10.1109/SP.2017.49>
- [13] Anirban Chakraborty, Manaar Alam, Vishal Dey, Anupam Chattopadhyay, and Debdeep Mukhopadhyay. 2021. A survey on adversarial attacks and defences. *CAAI Transactions on Intelligence Technology* 6, 1 (March 2021), 25–45. DOI : <https://doi.org/10.1049/cit2.12028>

- [14] Patrick Chao, Edoardo Debenedetti, Alexander Robey, Maksym Andriushchenko, Francesco Croce, Vikash Sehwal, Edgar Dobriban, Nicolas Flammarion, George J. Pappas, Florian Tramèr, et al. 2024. JailbreakBench: An open robustness benchmark for jailbreaking large language models. In *Proceedings of the 38th International Conference on Neural Information Processing Systems* (Vancouver, BC, Canada). Curran Associates Inc., Red Hook, NY, USA, Article 1745, 25 pages. Retrieved from https://proceedings.neurips.cc/paper_files/paper/2024/file/63092d79154adebd7305dfd498cbff70-Paper-Datasets_and_Benchmarks_Track.pdf
- [15] Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J. Pappas, and Eric Wong. 2025. Jailbreaking black box large language models in twenty queries. In *2025 IEEE Conference on Secure and Trustworthy Machine Learning*. IEEE, Copenhagen, Denmark, 23–42. DOI : <https://doi.org/10.1109/SaTML64287.2025.00010>
- [16] Aobo Chen, Yangyi Li, Chenxu Zhao, and Mengdi Huai. 2025. A survey of security and privacy issues of machine unlearning. *AI Magazine* 46, 1 (Jan. 2025), 12 pages. DOI : <https://doi.org/10.1002/aaai.12209>
- [17] Shuo Chen, Zhen Han, Bailan He, Zifeng Ding, Wenqian Yu, Philip Torr, Volker Tresp, and Jindong Gu. 2024. Red Teaming GPT-4V: Are GPT-4V Safe Against Uni/Multi-Modal Jailbreak Attacks? arXiv:2404.03411 [cs.LG].
- [18] Xinlei Chen, Hao Fang, Tsung-Yi Lin, Ramakrishna Vedantam, Saurabh Gupta, Piotr Dollár, and C. Lawrence Zitnick. 2015. Microsoft COCO Captions: Data Collection and Evaluation Server. arXiv:1504.00325 [cs.CV].
- [19] Nhat Chung, Sensen Gao, Tuan-Anh Vu, Jie Zhang, Aishan Liu, Yun Lin, Jin Song Dong, and Qing Guo. 2024. Towards Transferable Attacks Against Vision-LLMs in Autonomous Driving with Typography. arXiv:2405.14169 [cs.CV].
- [20] Mauro Conti, Nicola Dragoni, and Viktor Lesyk. 2016. A survey of man in the middle attacks. *IEEE Communications Surveys & Tutorials* 18, 3 (2016), 2027–2051. DOI : <https://doi.org/10.1109/COMST.2016.2548426>
- [21] Andrei Costin, Jonas Zaddach, Aurélien Francillon, and Davide Balzarotti. 2014. A large-scale analysis of the security of embedded firmwares. In *23rd USENIX Security Symposium*. USENIX Association, San Diego, CA, 95–110. Retrieved from <https://www.usenix.org/conference/usenixsecurity14/technical-sessions/presentation/costin>
- [22] Xuanming Cui, Alejandro Aparcedo, Young Kyun Jang, and Ser-Nam Lim. 2024. On the robustness of large multimodal models against image adversarial attacks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Seattle, WA, USA, 24625–24634. DOI : <https://doi.org/10.1109/CVPR52733.2024.02325>
- [23] Kelton A.P. da Costa, João P. Papa, Celso O. Lisboa, Roberto Munoz, and Victor Hugo C. de Albuquerque. 2019. Internet of things: A survey on machine learning-based intrusion detection approaches. *Computer Networks* 151 (2019), 147–157. DOI : <https://doi.org/10.1016/j.comnet.2019.01.023>
- [24] Gelei Deng, Yi Liu, Kailong Wang, Yuekang Li, Tianwei Zhang, and Yang Liu. 2024. Pandora: Jailbreak GPTs by Retrieval Augmented Generation Poisoning. arXiv:2402.08416 [cs.CR].
- [25] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. ImageNet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, Miami, FL, USA, 248–255. DOI : <https://doi.org/10.1109/CVPR.2009.5206848>
- [26] Yue Deng, Wenxuan Zhang, Sinno Jialin Pan, and Lidong Bing. 2023. Multilingual Jailbreak Challenges in Large Language Models. arXiv:2310.06474 [cs.CL].
- [27] Thierry Deruyttere, Simon Vandenhende, Dusan Grujicic, Luc Van Gool, and Marie-Francine Moens. 2019. Talk2Car: Taking control of your self-driving car. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*. Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (Eds.). Association for Computational Linguistics, Hong Kong, China, 2088–2098. DOI : <https://doi.org/10.18653/v1/D19-1215>
- [28] Lucas Dixon, John Li, Jeffrey Sorensen, Nithum Thain, and Lucy Vasserman. 2018. Measuring and mitigating unintended bias in text classification. In *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society* (New Orleans, LA, USA). Association for Computing Machinery, New York, NY, USA, 67–73. DOI : <https://doi.org/10.1145/3278721.3278729>
- [29] Yingkai Dong, Zheng Li, Xiangtao Meng, Ning Yu, and Shanqing Guo. 2024. Fuzz-Testing Meets LLM-Based Agents: An Automated and Efficient Framework for Jailbreaking Text-To-Image Generation Models. arXiv:2408.00523 [cs.CR].
- [30] Wenli Duo, MengChu Zhou, and Abdullah Abusorrah. 2022. A survey of cyber attacks on cyber physical systems: Recent advances and challenges. *IEEE/CAA Journal of Automatica Sinica* 9, 5 (2022), 784–800. DOI : <https://doi.org/10.1109/JAS.2022.105548>
- [31] European Commission. 2021. Proposal for a Regulation of the European Parliament and of the Council Laying Down Harmonised Rules on Artificial Intelligence (Artificial Intelligence Act) and Amending Certain Union Legislative Acts. Retrieved from <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52021PC0206>
- [32] European Parliament and Council of the European Union. 2016. Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation). Official Journal of the European Union, L 119, 1–88. Retrieved from <https://eur-lex.europa.eu/eli/reg/2016/679/oj>
- [33] Kevin Eykholt, Ivan Evtimov, Earlene Fernandes, Bo Li, Amir Rahmati, Chaowei Xiao, Atul Prakash, Tadayoshi Kohno, and Dawn Song. 2018. Robust physical-world attacks on deep learning visual classification. In *2018 IEEE/CVF*

- Conference on Computer Vision and Pattern Recognition*. IEEE, Salt Lake City, UT, USA, 1625–1634. DOI : <https://doi.org/10.1109/CVPR.2018.00175>
- [34] Hao-Shu Fang, Chenxi Wang, Minghao Gou, and Cewu Lu. 2020. GraspNet-1Billion: A large-scale benchmark for general object grasping. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Seattle, WA, USA, 11441–11450. DOI : <https://doi.org/10.1109/CVPR42600.2020.01146>
 - [35] Karine Gandolfi, Christophe Mourtel, and Francis Olivier. 2001. Electromagnetic analysis: Concrete results. In *Cryptographic Hardware and Embedded Systems — CHES 2001*. Springer Berlin, Berlin, 251–261. DOI : https://doi.org/10.1007/3-540-44709-1_21
 - [36] Kuofeng Gao, Yang Bai, Jiawang Bai, Yong Yang, and Shu-Tao Xia. 2024. Adversarial Robustness for Visual Grounding of Multimodal Large Language Models. arXiv:2405.09981 [cs.CV].
 - [37] Sensen Gao, Xiaojun Jia, Yihao Huang, Ranjie Duan, Jindong Gu, Yang Bai, Yang Liu, and Qing Guo. 2024. HTS-Attack: Heuristic Token Search for Jailbreaking Text-to-Image Models. arXiv:2408.13896 [cs.CV].
 - [38] Samuel Gehman, Suchin Gururangan, Maarten Sap, Yejin Choi, and Noah A. Smith. 2020. RealToxicityPrompts: Evaluating neural toxic degeneration in language models. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, Trevor Cohn, Yulan He, and Yang Liu (Eds.). Association for Computational Linguistics, Online, 3356–3369. DOI : <https://doi.org/10.18653/v1/2020.findings-emnlp.301>
 - [39] Yichen Gong, Delong Ran, Jinyuan Liu, Conglei Wang, Tianshuo Cong, Anyu Wang, Sisi Duan, and Xiaoyun Wang. 2025. FigStep: Jailbreaking large vision-language models via typographic visual prompts. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. AI Magazine, Washington, USA, 23951–23959. DOI : <https://doi.org/10.1609/aaai.v39i22.34568>
 - [40] Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. 2014. Explaining and Harnessing Adversarial Examples. arXiv:1412.6572 [stat.ML].
 - [41] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. 2017. Making the V in VQA matter: Elevating the role of image understanding in visual question answering. In *2017 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, Honolulu, HI, USA, 6325–6334. DOI : <https://doi.org/10.1109/CVPR.2017.670>
 - [42] Marc Gruber, Ian C. MacMillan, and James D. Thompson. 2008. Look before you leap: market opportunity identification in emerging technology firms. *Management Science* 54, 9 (2008), 1652–1665. DOI : <https://doi.org/10.1287/mnsc.1080.0877>
 - [43] Kevin Immanuel Gubbi, Banafsheh Saber Latibari, Anirudh Srikanth, Tyler Sheaves, Sayed Arash Beheshti-Shirazi, Sai Manoj P. D., Satareh Rafatirad, Avesta Sasan, Houman Homayoun, and Soheil Salehi. 2023. Hardware trojan detection using machine learning: A tutorial. *ACM Transactions on Embedded Computing Systems* 22, 3, Article 46 (April 2023), 26 pages. DOI : <https://doi.org/10.1145/3579823>
 - [44] Ujjwal Guin, Daniel DiMase, and Mohammad Tehranipoor. 2014. Counterfeit integrated circuits: Detection, avoidance, and the challenges ahead. *Journal of Electronic Testing* 30, 1 (February 2014), 9–23. DOI : <https://doi.org/10.1007/s10836-013-5430-8>
 - [45] Kazuki Hayashi, Kazuma Onishi, Toma Suzuki, Yusuke Ide, Seiji Gobara, Shigeki Saito, Yusuke Sakai, Hidetaka Kamigaito, Katsuhiko Hayashi, and Taro Watanabe. 2025. IRR: Image review ranking framework for evaluating vision-language models. In *Proceedings of the 31st International Conference on Computational Linguistics*. Owen Rambow, Leo Wanner, Marianna Apidianaki, Hend Al-Khalifa, Barbara Di Eugenio, and Steven Schockaert (Eds.). Association for Computational Linguistics, Abu Dhabi, UAE, 9939–9956. <https://aclanthology.org/2025.coling-main.665/>
 - [46] Xuanli He, Jun Wang, Qiongkai Xu, Pasquale Minervini, Pontus Stenetorp, Benjamin I. P. Rubinstein, and Trevor Cohn. 2025. TUBA: Cross-lingual transferability of backdoor attacks in LLMs with instruction tuning. In *Findings of the Association for Computational Linguistics: ACL 2025*. Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar (Eds.). Association for Computational Linguistics, Vienna, Austria, 16504–16544. DOI : <https://doi.org/10.18653/v1/2025.findings-acl.848>
 - [47] Xiaowei Huang, Daniel Kroening, Wenjie Ruan, James Sharp, Youcheng Sun, Emese Thamo, Min Wu, and Xinping Yi. 2020. A survey of safety and trustworthiness of deep neural networks: Verification, testing, adversarial attack and defence, and interpretability. *Computer Science Review* 37 (2020), 100270. DOI : <https://doi.org/10.1016/j.cosrev.2020.100270>
 - [48] Yiyang Huang, Zixuan Wang, Zishen Wan, Yapeng Tian, Haobo Xu, Yinhe Han, and Yiming Gan. 2025. ANNIE: Be Careful of Your Robots. arXiv:2509.03383 [cs.AI].
 - [49] Abdulmalik Humayed, Jingqiang Lin, Fengjun Li, and Bo Luo. 2017. Cyber-physical systems security—a survey. *IEEE Internet of Things Journal* 4, 6 (2017), 1802–1831. DOI : <https://doi.org/10.1109/JIOT.2017.2703172>
 - [50] Todd E. Humphreys, Jahshan Bhatti, Daniel Shepard, and Kyle Wesson. 2012. The Texas spoofing test battery: Toward a standard for evaluating GPS signal authentication techniques. In *Proceedings of the 25th International Technical Meeting of the Satellite Division of The Institute of Navigation*. The Institute of Navigation, Nashville, TN, 3569–3583. Retrieved from <https://www.ion.org/publications/abstract.cfm?articleID=10532>

- [51] Todd E. Humphreys, Brent M. Ledvina, Mark L. Psiaki, Brady W. O'Hanlon, and Paul M. Kintner, Jr. 2008. Assessing the spoofing threat: Development of a portable GPS civilian spoofer. In *Proceedings of the 21st International Technical Meeting of the Satellite Division of The Institute of Navigation*. The Institute of Navigation, Savannah, GA, 2314–2325. Retrieved from <https://www.ion.org/publications/abstract.cfm?articleID=8132>
- [52] Hakan Inan, Khanh Upasani, Jianfeng Chi, Rashmi Rungta, Krithika Iyer, Yuning Mao, Michael Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madian Khabza. 2023. Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations. arXiv:2312.06674 [cs.CL].
- [53] International Organization for Standardization. 2023. Safety of Machinery — Safety-related Parts of Control Systems — Part 1: General Principles for Design. Retrieved from <https://www.iso.org/standard/73481.html>
- [54] International Organization for Standardization. 2025. Robotics — Safety requirements — Part 1: Industrial robots. Retrieved from <https://www.iso.org/standard/73933.html>
- [55] Jiaming Ji, Mickel Liu, Josef Dai, Xuehai Pan, Chi Zhang, Ce Bian, Boyuan Chen, Ruiyang Sun, Yizhou Wang, and Yaodong Yang. 2023. BeaverTails: Towards improved safety alignment of LLM via a human-preference dataset. In *Advances in Neural Information Processing Systems*, A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine (Eds.), Vol. 36. Curran Associates, Inc., Vancouver, Canada, 24678–24704. Retrieved from https://proceedings.neurips.cc/paper_files/paper/2023/file/4dbb61cb68671edc4ca3712d70083b9f-Paper-Datasets_and_Benchmarks.pdf
- [56] Daniel Kang, Xuechen Li, Ion Stoica, Carlos Guestrin, Matei Zaharia, and Tatsunori Hashimoto. 2024. Exploiting programmatic behavior of LLMs: Dual-use through standard security attacks. In *2024 IEEE Security and Privacy Workshops*. IEEE, San Francisco, California, USA, 132–143. DOI: <https://doi.org/10.1109/SPW63631.2024.00018>
- [57] Sahar Kazemzadeh, Vicente Ordonez, Mark Matten, and Tamara Berg. 2014. ReferItGame: Referring to objects in photographs of natural scenes. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Alessandro Moschitti, Bo Pang, and Walter Daelemans (Eds.). Association for Computational Linguistics, Doha, Qatar, 787–798. DOI: <https://doi.org/10.3115/v1/D14-1086>
- [58] Hyungsub Kim, Riwat Bandyopadhyay, Muslum Ozgur Ozmen, Z. Berkay Celik, Antonio Bianchi, Yongdae Kim, and Dongyan Xu. 2024. A systematic study of physical sensor attack hardness. In *2024 IEEE Symposium on Security and Privacy*. IEEE, San Francisco, California, USA, 2328–2347. DOI: <https://doi.org/10.1109/SP54263.2024.00143>
- [59] Moo Jin Kim, Karl Pertsch, Siddharth Karamcheti, Ted Xiao, Ashwin Balakrishna, Suraj Nair, Rafael Rafailov, Ethan P. Foster, Pannag R. Sanketi, Quan Vuong, et al. 2025. OpenVLA: An open-source vision-language-action model. In *Proceedings of The 8th Conference on Robot Learning*, Pulkit Agrawal, Oliver Kroemer, and Wolfram Burgard (Eds.), Vol. 270. PMLR, Munich, Germany, 2679–2713. Retrieved from <https://proceedings.mlr.press/v270/kim25c.html>
- [60] Won Tae Kim, Won Sik Ham, Wooju Jeong, Hyun Jung Song, Koon Ho Rha, and Young Deuk Choi. 2009. Failure and malfunction of da vinci surgical systems during various robotic surgeries: Experience from six departments at a single institute. *Urology* 74, 6 (2009), 1234–1237. DOI: <https://doi.org/10.1016/j.urology.2009.05.071>
- [61] Paul Kocher, Joshua Jaffe, and Benjamin Jun. 1999. Differential Power Analysis. In *Advances in Cryptology — CRYPTO'99*, Michael Wiener (Ed.). Springer Berlin Heidelberg, Berlin, 388–397. DOI: https://doi.org/10.1007/3-540-48405-1_25
- [62] Haoran Li, Dadi Guo, Wei Fan, Mingshi Xu, Jie Huang, Fanpu Meng, and Yangqiu Song. 2023. Multi-step jailbreaking privacy attacks on chatGPT. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 4138–4153. DOI: <https://doi.org/10.18653/v1/2023.findings-emnlp.272>
- [63] Mukai Li, Lei Li, Yuwei Yin, Masood Ahmed, Zhenguang Liu, and Qi Liu. 2024. Red teaming visual language models. In *Findings of the Association for Computational Linguistics: ACL 2024*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 3326–3342. DOI: <https://doi.org/10.18653/v1/2024.findings-acl.198>
- [64] Shuai Li, Chen Yizhe, Li Dong, Liu Sichao, Lan Dapeng, Liu Yu, and Zhibo Pang. 2025. Transferring Vision-Language-Action Models to Industry Applications: Architectures, Performance, and Challenges. arXiv:2509.23121 [cs.AI].
- [65] Tung-Ling Li and Hongliang Liu. 2025. Logit-Gap Steering: Efficient Short-Suffix Jailbreaks for Aligned Large Language Models. arXiv:2506.24056 [cs.CR].
- [66] Xuan Li, Zhanke Zhou, Jianing Zhu, Jiangchao Yao, Tongliang Liu, and Bo Han. 2023. DeepInception: Hypnotize Large Language Model to Be Jailbreaker. arXiv:2311.03191 [cs.CL].
- [67] Yifan Li, Hangyu Guo, Kun Zhou, Wayne Xin Zhao, and Ji-Rong Wen. 2025. Images are achilles' heel of alignment: Exploiting visual vulnerabilities for jailbreaking multimodal large language models. In *Computer Vision – ECCV 2024*. Ales Leonardis, Elisa Ricci, Stefan Roth, Olga Russakovsky, Torsten Sattler, and Gül Varol (Eds.). Springer Nature Switzerland, Cham, 174–189. DOI: https://doi.org/10.1007/978-3-031-73464-9_11
- [68] Yuxi Li, Yi Liu, Yuekang Li, Ling Shi, Gelei Deng, Shengquan Chen, and Kailong Wang. 2025. Lockpicking LLMs: A Logit-Based Jailbreak Using Token-level Manipulation. arXiv:2405.13068 [cs.CR].
- [69] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. 2023. Code as policies: Language model programs for embodied control. In *2023 IEEE International Conference on Robotics and Automation*. IEEE, London, United Kingdom, 9493–9500. DOI: <https://doi.org/10.1109/ICRA48891.2023.10160591>

- [70] Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 3214–3252. DOI : <https://doi.org/10.18653/v1/2022.acl-long.229>
- [71] Zi Lin, Zihan Wang, Yongqi Tong, Yangkun Wang, Yuxin Guo, Yujia Wang, and Jingbo Shang. 2023. ToxicChat: Unveiling hidden challenges of toxicity detection in real-world user-AI conversation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 4694–4702. DOI : <https://doi.org/10.18653/v1/2023.findings-emnlp.311>
- [72] Daizong Liu, Mingyu Yang, Xiaoye Qu, Pan Zhou, Yu Cheng, and Wei Hu. 2025. A survey of attacks on large vision–language models: Resources, advances, and future trends. *IEEE Transactions on Neural Networks and Learning Systems* 36, 11 (2025), 19525–19545. DOI : <https://doi.org/10.1109/TNNLS.2025.3592935>
- [73] Hongyi Liu, Shaochen Zhong, Xintong Sun, Minghao Tian, Mohsen Hariri, Zirui Liu, Ruixiang Tang, Zhimeng Jiang, Jiayi Yuan, Yu-Neng Chuang, et al. 2025. LoRATK: LoRA once, backdoor everywhere in the share-and-play ecosystem. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng (Eds.). Association for Computational Linguistics, Suzhou, China, 23009–23047. DOI : <https://doi.org/10.18653/v1/2025.findings-emnlp.1253>
- [74] Xin Liu, Yichen Zhu, Jindong Gu, Yunshi Lan, Chao Yang, and Yu Qiao. 2025. MM-SafetyBench: A benchmark for safety evaluation of multimodal large language models. In *Computer Vision – ECCV 2024*, Aleš Leonardis, Elisa Ricci, Stefan Roth, Olga Russakovsky, Torsten Sattler, and Gül Varol (Eds.). Springer Nature Switzerland, Cham, 386–403. DOI : https://doi.org/10.1007/978-3-031-72992-8_22
- [75] Stefano Longari, Jacopo Jannone, Mario Polino, Michele Carminati, Andrea Zanchettin, Mara Tanelli, and Stefano Zanero. 2025. Janus: A trusted execution environment approach for attack detection in industrial robot controllers. *IEEE Transactions on Emerging Topics in Computing* 13, 1 (2025), 185–195. DOI : <https://doi.org/10.1109/TETC.2024.3390435>
- [76] Tim Loughran, Bill McDonald, and Hayong Yun. 2008. A wolf in sheep’s clothing: The use of ethics-related terms in 10-K reports. *Journal of Business Ethics* 89, 1 (May 2008), 39–49. DOI : <https://doi.org/10.1007/s10551-008-9910-1>
- [77] Weidi Luo, Siyuan Ma, Xiaogeng Liu, Xiaoyu Guo, and Chaowei Xiao. 2024. JailBreakV: A Benchmark for Assessing the Robustness of MultiModal Large Language Models against Jailbreak Attacks. arXiv:2404.03027 [cs.CR].
- [78] Huijie Lv, Xiao Wang, Yuan Zhang, Caishuang Huang, Shihan Dou, Junjie Ye, Tao Gui, Qi Zhang, and Xuan-jing Huang. 2024. CodeChameleon: Personalized Encryption Framework for Jailbreaking Large Language Models. arXiv:2402.16717 [cs.CL].
- [79] Jiachen Ma, Yijiang Li, Zhiqing Xiao, Anda Cao, Jie Zhang, Chao Ye, and Junbo Zhao. 2025. Jailbreaking prompt attack: A controllable adversarial attack against diffusion models. In *Findings of the Association for Computational Linguistics: NAACL 2025*, Luis Chiruzzo, Alan Ritter, and Lu Wang (Eds.). Association for Computational Linguistics, Albuquerque, New Mexico, 3141–3157. DOI : <https://doi.org/10.18653/v1/2025.findings-naacl.172>
- [80] Siyuan Ma, Weidi Luo, Yu Wang, Xiaogeng Liu, Muhao Chen, Bo Li, and Chaowei Xiao. 2024. Visual-RolePlay: Universal Jailbreak Attack on MultiModal Large Language Models via Role-playing Image Character. arXiv:2405.20773 [cs.CR].
- [81] Yanmao Man, Ming Li, and Ryan Gerdes. 2020. GhostImage: Remote perception attacks against camera-based image classification systems. In *23rd International Symposium on Research in Attacks, Intrusions and Defenses*. USENIX Association, San Sebastian, 317–332. Retrieved from <https://www.usenix.org/conference/raid2020/presentation/man>
- [82] Junhua Mao, Jonathan Huang, Alexander Toshev, Oana Camburu, Alan Yuille, and Kevin Murphy. 2016. Generation and comprehension of unambiguous object descriptions. In *2016 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, Las Vegas, NV, USA, 11–20. DOI : <https://doi.org/10.1109/CVPR.2016.9>
- [83] Kenneth Marino, Mohammad Rastegari, Ali Farhadi, and Roozbeh Mottaghi. 2019. OK-VQA: A visual question answering benchmark requiring external knowledge. In *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Long Beach, CA, USA, 3190–3199. DOI : <https://doi.org/10.1109/CVPR.2019.00331>
- [84] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. 2024. HarmBench: A standardized evaluation framework for automated red teaming and robust refusal. In *Proceedings of the 41st International Conference on Machine Learning*, Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (Eds.), Vol. 235. PMLR, Vienna, Austria, 35181–35224. Retrieved from <https://proceedings.mlr.press/v235/mazeika24a.html>
- [85] Tim McBride, Anne Townsend, Michael Ekstrom, Lauren Lusty, and Julian Sexton. 2018. Data integrity: Recovering from ransomware and other destructive events. In *2018 IEEE Cybersecurity Development*. IEEE, Cambridge, MA, USA, 140–140. DOI : <https://doi.org/10.1109/SecDev.2018.00036>
- [86] Noriyuki Miura, Tatsuya Machida, Kohei Matsuda, Makoto Nagata, Shoei Nashimoto, and Daisuke Suzuki. 2019. A low-cost replica-based distance-spoofing attack on mmwave FMCW radar. In *Proceedings of the 3rd ACM Workshop on Attacks and Solutions in Hardware Security Workshop* (London, United Kingdom). Association for Computing Machinery, New York, NY, USA, 95–100. DOI : <https://doi.org/10.1145/3338508.3359567>

- [87] Moin Nadeem, Anna Bethke, and Siva Reddy. 2021. StereoSet: Measuring stereotypical bias in pretrained language models. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, Chengqing Zong, Fei Xia, Wenjie Li, and Roberto Navigli (Eds.). Association for Computational Linguistics, Online, 5356–5371. DOI: <https://doi.org/10.18653/v1/2021.acl-long.416>
- [88] Ben Nassi, Yisroel Mirsky, Dudi Nassi, Raz Ben-Netanel, Oleg Drokin, and Yuval Elovici. 2020. Phantom of the ADAS: Securing advanced driver-assistance systems from split-second phantom attacks. In *Proceedings of the 2020 ACM SIGSAC Conference on Computer and Communications Security* (Virtual Event, USA). Association for Computing Machinery, New York, NY, USA, 293–308. DOI: <https://doi.org/10.1145/3372297.3423359>
- [89] Subash Neupane, Shaswata Mitra, Ivan A. Fernandez, Swayamjit Saha, Sudip Mittal, Jingdao Chen, Nisha Pillai, and Shahram Rahimi. 2024. Security considerations in AI-robotics: A survey of current methods, challenges, and opportunities. *IEEE Access* 12 (2024), 22072–22097. DOI: <https://doi.org/10.1109/ACCESS.2024.3363657>
- [90] Anh Nguyen, Jason Yosinski, and Jeff Clune. 2015. Deep neural networks are easily fooled: High confidence predictions for unrecognizable images. In *2015 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, Boston, MA, USA, 427–436. DOI: <https://doi.org/10.1109/CVPR.2015.7298640>
- [91] Zhenxing Niu, Haodong Ren, Xinbo Gao, Gang Hua, and Rong Jin. 2024. Jailbreaking Attack against Multimodal Large Language Model. arXiv:2402.02309 [cs.LG].
- [92] Nicolas Papernot, Patrick McDaniel, and Ian J. Goodfellow. 2016. Transferability in Machine Learning: from Phenomena to Black-Box Attacks using Adversarial Samples. arXiv:1605.07277 [cs.CR].
- [93] Jonathan Petit, Bas Stottelaar, Michael Feiri, and Frank Kargl. 2015. Remote Attacks on Automated Vehicles Sensors: Experiments on Camera and LiDAR. Black Hat Europe. Retrieved from <https://www.blackhat.com/docs/eu-15/materials/eu-15-Petit-Self-Driving-And-Connected-Cars-Fooling-Sensors-And-Tracking-Drivers-wp1.pdf>
- [94] Bryan A. Plummer, Liwei Wang, Chris M. Cervantes, Juan C. Caicedo, Julia Hockenmaier, and Svetlana Lazebnik. 2015. Flickr30k entities: Collecting region-to-phrase correspondences for richer image-to-sentence models. In *2015 IEEE International Conference on Computer Vision*. IEEE, Santiago, Chile, 2641–2649. DOI: <https://doi.org/10.1109/ICCV.2015.303>
- [95] Mark L. Psiaki and Todd E. Humphreys. 2016. GNSS spoofing and detection. *Proc. IEEE* 104, 6 (2016), 1258–1270. DOI: <https://doi.org/10.1109/JPROC.2016.2526658>
- [96] Xiangyu Qi, Yi Zeng, Tinghao Xie, Pin-Yu Chen, Ruoxi Jia, Prateek Mittal, and Peter Henderson. 2023. Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To! arXiv:2310.03693 [cs.CL].
- [97] Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. 2023. Direct preference optimization: Your language model is secretly a reward model. In *Advances in Neural Information Processing Systems*, A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine (Eds.), Vol. 36. Curran Associates, Inc., New Orleans, Louisiana, USA, 53728–53741. Retrieved from https://proceedings.neurips.cc/paper_files/paper/2023/file/a85b405ed65c6477afe8302b5e06ce7-Paper-Conference.pdf
- [98] Rujit Raval, Alison Maskus, Benjamin Saltmiras, Morgan Dunn, Peter J. Hawrylak, and John Hale. 2018. Competitive learning environment for cyber-physical system security experimentation. In *2018 1st International Conference on Data Intelligence and Security*. IEEE, South Padre Island, TX, USA, 211–218. DOI: <https://doi.org/10.1109/ICDIS.2018.00042>
- [99] Delphine Reinhardt and Archan Misra. 2025. Toward privacy-respecting service robots. *IEEE Pervasive Computing* 24, 3 (2025), 63–72. DOI: <https://doi.org/10.1109/MPRV.2025.3555407>
- [100] Ishtiaq Rouf, Rob Miller, Hossen Mustafa, Travis Taylor, Sangho Oh, Wenyuan Xu, Marco Gruteser, Wade Trappe, and Ivan Seskar. 2010. Security and privacy vulnerabilities of in-car wireless networks: A tire pressure monitoring system case study. In *19th USENIX Security Symposium*. USENIX Association, Washington, DC, USA, 323–338. Retrieved from <https://dl.acm.org/doi/10.5555/1929820.1929848>
- [101] Rylan Schaeffer, Dan Valentine, Luke Bailey, James Chua, Cristóbal Eyzaguirre, Zane Durante, Joe Benton, Brando Miranda, Henry Sleight, John Hughes, et al. 2024. Failures to Find Transferable Image Jailbreaks Between Vision-Language Models. arXiv:2407.15211 [cs.CV].
- [102] Wenqi Shao, Meng Lei, Yutao Hu, Peng Gao, Peng Xu, Kaipeng Zhang, Fanqing Meng, Siyuan Huang, Hongsheng Li, Yu Qiao, and Ping Luo. 2025. TinyLVM-eHub: Towards Comprehensive and Efficient Evaluation for Large Vision-Language Models. *IEEE Transactions on Big Data* 11, 3 (2025), 933–947. DOI: <https://doi.org/10.1109/TBDATA.2025.3536930>
- [103] Junjie Shen, Jun Yeon Won, Zeyuan Chen, and Qi Alfred Chen. 2020. Drift with Devil: Security of Multi-Sensor Fusion based Localization in High-Level Autonomous Driving under GPS Spoofing. In *29th USENIX Security Symposium*. USENIX Association, Boston, MA, USA, 931–948. Retrieved from <https://www.usenix.org/conference/usenixsecurity20/presentation/shen>
- [104] Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. 2024. “Do anything now”: Characterizing and evaluating in-the-wild jailbreak prompts on large language models. In *Proceedings of the 2024 on ACM SIGSAC*

- Conference on Computer and Communications Security* (Salt Lake City, UT, USA). Association for Computing Machinery, New York, NY, USA, 1671–1685. DOI : <https://doi.org/10.1145/3658644.3670388>
- [105] Xinyue Shen, Yixin Wu, Michael Backes, and Yang Zhang. 2024. Voice Jailbreak Attacks Against GPT-4o. arXiv:2405.19103 [cs.CR].
 - [106] Emily Sheng, Kai-Wei Chang, Premkumar Natarajan, and Nanyun Peng. 2019. The woman worked as a babysitter: On biases in language generation. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan (Eds.). Association for Computational Linguistics, Hong Kong, China, 3407–3412. DOI : <https://doi.org/10.18653/v1/D19-1339>
 - [107] Junli Shi, Yuan Dai, Yu Cheng, Sai Xie, Gang Li, Yuan Liu, Jingxiao Wang, Ruirui Zhang, Ningning Bai, Minkun Cai, et al. 2023. Embedment of sensing elements for robust, highly sensitive, and cross-talk-free iontronic skins for robotics applications. *Science Advances* 9, 9 (2023), ead8831. DOI : <https://doi.org/10.1126/sciadv.adf8831>
 - [108] Mohit Shridhar, Lucas Manuelli, and Dieter Fox. 2023. Perceiver-actor: A multi-task transformer for robotic manipulation. In *Proceedings of The 6th Conference on Robot Learning*, Karen Liu, Dana Kulic, and Jeff Ichnowski (Eds.), Vol. 205. PMLR, Auckland, New Zealand, 785–799. Retrieved from <https://proceedings.mlr.press/v205/shridhar23a.html>
 - [109] Mohit Shridhar, Jesse Thomason, Daniel Gordon, Yonatan Bisk, Winson Han, Roozbeh Mottaghi, Luke Zettlemoyer, and Dieter Fox. 2020. ALFRED: A benchmark for interpreting grounded instructions for everyday tasks. In *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Virtual (originally planned for Seattle, WA, USA), 10737–10746. DOI : <https://doi.org/10.1109/CVPR42600.2020.01075>
 - [110] Derek J. Straub. 2020. Detecting the presence of fog using low-cost proximity sensors. *Aerosol and Air Quality Research* 20, 5 (May 2020), 981–990. DOI : <https://doi.org/10.4209/aaqr.2019.11.0612>
 - [111] Hao Sun, Zhixin Zhang, Jiawen Deng, Jiale Cheng, and Minlie Huang. 2023. Safety Assessment of Chinese Large Language Models. arXiv:2304.10436 [cs.CL].
 - [112] Jiachen Sun, Yulong Cao, Qi Alfred Chen, and Z. Morley Mao. 2020. Towards robust LiDAR-based perception in autonomous driving: General black-box adversarial sensor attack and countermeasures. In *29th USENIX Security Symposium*. USENIX Association, Boston, MA, USA, 877–894. Retrieved from <https://www.usenix.org/conference/usenixsecurity20/presentation/sun>
 - [113] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. 2013. Intriguing Properties of Neural Networks. arXiv:1312.6199 [cs.CV].
 - [114] Timothy Trippel, Ofir Weisse, Wenyan Xu, Peter Honeyman, and Kevin Fu. 2017. WALNUT: Waging doubt on the integrity of MEMS accelerometers with acoustic injection attacks. In *2017 IEEE European Symposium on Security and Privacy*. IEEE, Paris, France, 3–18. DOI : <https://doi.org/10.1109/EuroSP.2017.42>
 - [115] Haoqin Tu, Chenhang Cui, Zijun Wang, Yiyang Zhou, Bingchen Zhao, Junlin Han, Wangchunshu Zhou, Huaxiu Yao, and Cihang Xie. 2023. How Many Unicorns are in This Image? A Safety Evaluation Benchmark for Vision LLMs. arXiv:2311.16101 [cs.CV].
 - [116] Yazhou Tu, Sara Rampazzi, and Xiali Hei. 2022. Towards Adversarial Control Loops in Sensor Attacks: A Case Study to Control the Kinematics and Actuation of Embedded Systems. arXiv:2203.07670 [cs.CR].
 - [117] Hao Wang, Hao Li, Minlie Huang, and Lei Sha. 2024. ASETF: A novel method for jailbreak attack on LLMs through translate suffix embeddings. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen (Eds.). Association for Computational Linguistics, Miami, Florida, USA, 2697–2711. DOI : <https://doi.org/10.18653/v1/2024.emnlp-main.157>
 - [118] Xingyao Wang, Yangyi Chen, Lifan Yuan, Yizhe Zhang, Yunzhu Li, Hao Peng, and Heng Ji. 2024. Executable code actions elicit better LLM agents. In *Proceedings of the 41st International Conference on Machine Learning*, Ruslan Salakhutdinov, Zico Kolter, Katherine Heller, Adrian Weller, Nuria Oliver, Jonathan Scarlett, and Felix Berkenkamp (Eds.), Vol. 235. PMLR, Vienna, Austria, 50208–50232. Retrieved from <https://proceedings.mlr.press/v235/wang24h.html>
 - [119] Xinpeng Wang, Xiaoyuan Yi, Han Jiang, Shanlin Zhou, Zhihua Wei, and Xing Xie. 2023. ToViLaG: Your visual-language generative model is also an evildoer. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 3508–3533. DOI : <https://doi.org/10.18653/v1/2023.emnlp-main.213>
 - [120] Zeming Wei, Yifei Wang, Ang Li, Yichuan Mo, and Yisen Wang. 2023. Jailbreak and Guard Aligned Language Models with Only Few In-Context Demonstrations. arXiv:2310.06387 [cs.LG].
 - [121] Wikipedia. 2024. List of Tesla Autopilot crashes. Retrieved from https://en.wikipedia.org/wiki/List_of_Tesla_Autopilot_crashes. Accessed: 2024-11-04.
 - [122] Peng Xu, Wenqi Shao, Kaipeng Zhang, Peng Gao, Shuo Liu, Meng Lei, Fanqing Meng, Siyuan Huang, Yu Qiao, and Ping Luo. 2025. LVLMEHUB: A comprehensive evaluation benchmark for large vision-language models. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47, 3 (2025), 1877–1893. DOI : <https://doi.org/10.1109/TPAMI.2024.3507000>

- [123] Yuan Xu, Xingshuo Han, Gelei Deng, Jiwei Li, Yang Liu, and Tianwei Zhang. 2023. SoK: Rethinking sensor spoofing attacks against robotic vehicles from a systematic view. In *2023 IEEE 8th European Symposium on Security and Privacy*. IEEE, Delft, Netherlands, 1082–1100. DOI: <https://doi.org/10.1109/EuroSP57164.2023.00067>
- [124] Xianjun Yang, Xiao Wang, Qi Zhang, Linda Petzold, William Yang Wang, Xun Zhao, and Dahua Lin. 2023. Shadow Alignment: The Ease of Subverting Safely-Aligned Language Models. arXiv:2310.02949 [cs.CL].
- [125] De Jong Yeong, Gustavo Velasco-Hernandez, John Barry, and Joseph Walsh. 2021. Sensor and sensor fusion technology in autonomous vehicles: A review. *Sensors* 21, 6 (2021). DOI: <https://doi.org/10.3390/s21062140>
- [126] Zonghao Ying, Aishan Liu, Xianglong Liu, and Dacheng Tao. 2024. Unveiling the Safety of GPT-4o: An Empirical Study using Jailbreak Attacks. arXiv:2406.06302 [cs.CR].
- [127] Zheng-Xin Yong, Cristina Menghini, and Stephen H. Bach. 2023. Low-Resource Languages Jailbreak GPT-4. arXiv:2310.02446 [cs.CL].
- [128] Licheng Yu, Patrick Poirson, Shan Yang, Alexander C. Berg, and Tamara L. Berg. 2016. Modeling context in referring expressions. In *Computer Vision – ECCV 2016*, Bastian Leibe, Jiri Matas, Nicu Sebe, and Max Welling (Eds.). Springer International Publishing, Cham, 69–85. DOI: https://doi.org/10.1007/978-3-319-46475-6_5
- [129] Tianhe Yu, Deirdre Quillen, Zhanpeng He, Ryan Julian, Karol Hausman, Chelsea Finn, and Sergey Levine. 2020. Meta-World: A Benchmark and Evaluation for Multi-Task and Meta Reinforcement Learning. In *Proceedings of the Conference on Robot Learning*, Leslie Pack Kaelbling, Danica Kragic, and Komei Sugiura (Eds.), Vol. 100. PMLR, Virtual Event / Cambridge, MA, USA, 1094–1100. Retrieved from <https://proceedings.mlr.press/v100/yu20a.html>
- [130] Xuejing Yuan, Yuxuan Chen, Yue Zhao, Yunhui Long, Xiaokang Liu, Kai Chen, Shengzhi Zhang, Heqing Huang, XiaoFeng Wang, and Carl A. Gunter. 2018. CommanderSong: A systematic approach for practical adversarial voice recognition. In *27th USENIX Security Symposium*. USENIX Association, Baltimore, MD, 49–64. Retrieved from <https://www.usenix.org/conference/usenixsecurity18/presentation/yuan-xuejing>
- [131] Youliang Yuan, Wenxiang Jiao, Wenxuan Wang, Jen tse Huang, Pinjia He, Shuming Shi, and Zhaopeng Tu. 2024. GPT-4 Is Too Smart To Be Safe: Stealthy Chat with LLMs via Cipher. arXiv:2308.06463 [cs.CL].
- [132] Kexiong (Curtis) Zeng, Shinan Liu, Yuanchao Shu, Dong Wang, Haoyu Li, Yanzhi Dou, Gang Wang, and Yaling Yang. 2018. All your GPS Are belong to us: Towards stealthy manipulation of road navigation systems. In *27th USENIX Security Symposium*. USENIX Association, Baltimore, MD, 1527–1544. Retrieved from <https://www.usenix.org/conference/usenixsecurity18/presentation/zeng>
- [133] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. 2024. How johnny can persuade LLMs to jailbreak them: Rethinking persuasion to challenge AI safety by humanizing LLMs. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Lun-Wei Ku, Andre Martins, and Vivek Srikumar (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 14322–14350. DOI: <https://doi.org/10.18653/v1/2024.acl-long.773>
- [134] Guoming Zhang, Chen Yan, Xiaoyu Ji, Tianchen Zhang, Taimin Zhang, and Wenyan Xu. 2017. DolphinAttack: Inaudible voice commands. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security* (Dallas, Texas, USA). Association for Computing Machinery, New York, NY, USA, 103–117. DOI: <https://doi.org/10.1145/3133956.3134052>
- [135] Hao Zhang, Wenqi Shao, Hong Liu, Yongqiang Ma, Ping Luo, Yu Qiao, Nanning Zheng, and Kaipeng Zhang. 2025. B-AVIBench: Toward evaluating the robustness of large vision-language model on black-box adversarial visual-instructions. *IEEE Transactions on Information Forensics and Security* 20 (2025), 1434–1446. DOI: <https://doi.org/10.1109/TIFS.2024.3520306>
- [136] Kaiqi Zhang, Shuai Yuan, and Honghan Zhao. 2024. TALEC: Teach Your LLM to Evaluate in Specific Domain with In-house Criteria by Criteria Division and Zero-shot Plus Few-shot. arXiv:2407.10999 [cs.CL].
- [137] Yongting Zhang, Lu Chen, Guodong Zheng, Yifeng Gao, Rui Zheng, Jinlan Fu, Zhenfei Yin, Senjie Jin, Yu Qiao, Xuanjing Huang, et al. 2025. SPA-VL: A comprehensive safety preference alignment dataset for vision language models. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Nashville, TN, USA, 19867–19878. DOI: <https://doi.org/10.1109/CVPR52734.2025.01850>
- [138] Zhixin Zhang, Leqi Lei, Lindong Wu, Rui Sun, Yongkang Huang, Chong Long, Xiao Liu, Xuanyu Lei, Jie Tang, and Minlie Huang. 2023. SafetyBench: Evaluating the Safety of Large Language Models with Multiple Choice Questions. arXiv:2309.07045 [cs.CL]. DOI: <https://doi.org/10.48550/arXiv.2309.07045>
- [139] Zhixin Zhang, Junxiao Yang, Pei Ke, and Minlie Huang. 2023. Defending Large Language Models Against Jailbreaking Attacks Through Goal Prioritization. arXiv:2311.09096 [cs.CL]. DOI: <https://doi.org/10.48550/arXiv.2311.09096>
- [140] Lifeng Zhou and Vijay Kumar. 2023. Robust multi-robot active target tracking against sensing and communication attacks. *IEEE Transactions on Robotics* 39, 3 (2023), 1768–1780. DOI: <https://doi.org/10.1109/TRO.2022.3233341>
- [141] Man Zhou, Zhan Qin, Xiu Lin, Shengshan Hu, Qian Wang, and Kui Ren. 2019. Hidden voice commands: Attacks and defenses on the VCS of autonomous driving cars. *IEEE Wireless Communications* 26, 5 (2019), 128–133. DOI: <https://doi.org/10.1109/MWC.2019.1800477>

- [142] Sicheng Zhu, Ruiyi Zhang, Bang An, Gang Wu, Joe Barrow, Zichao Wang, Furong Huang, Ani Nenkova, and Tong Sun. 2023. AutoDAN: Interpretable Gradient-Based Adversarial Attacks on Large Language Models. arXiv:[2310.15140](#) [cs.CR].
- [143] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. 2023. Universal and Transferable Adversarial Attacks on Aligned Language Models. arXiv:[2307.15043](#) [cs.LG].

Received 9 June 2025; revised 3 February 2026; accepted 24 March 2026